

¹ Gain, not concomitant changes in spatial receptive field properties, improves task performance in a neural network attention model

⁵ Kai J Fox^{1*†}, Daniel Birman^{1*‡§}, Justin L Gardner¹

*For correspondence:

kaifox@stanford.edu (KF);
dbirman@uw.edu (DB)

[†]These authors contributed equally to this work

⁸ Present address: [§]Department of Biological Structure, University of Washington, USA

⁹ **Abstract** Attention allows us to focus sensory processing on behaviorally relevant aspects of the visual world. One potential mechanism of attention is a change in the gain of sensory responses. However, changing gain at early stages could have multiple downstream consequences for visual processing. Which, if any, of these effects can account for the benefits of attention for detection and discrimination? Using a model of primate visual cortex we document how a Gaussian-shaped gain modulation results in changes to spatial tuning properties. Forcing the model to use only these changes failed to produce any benefit in task performance. Instead, we found that gain alone was both necessary and sufficient to explain category detection and discrimination during attention. Our results show how gain can give rise to changes in receptive fields which are not necessary for enhancing task performance.

¹⁹ Introduction

²⁰ Deploying goal-directed spatial attention towards visual locations allows observers to detect targets with higher accuracy (*Hawkins et al., 1990*), faster reaction times (*Posner, 1980*), and higher sensitivity (*Sagi and Julesz, 1986*) providing humans and non-human primates with a mechanism to select and prioritize spatial visual information (*Carrasco, 2011*). These enhanced behavioral responses are accompanied by both an increase in the gain of sensory responses near attended locations (*Connor et al., 1996; McAdams and Maunsell, 1999*) and changes in the shape and size of receptive fields, typically shrinking and shifting towards the target of attention (*Ben Hamed et al., 2002; Womelsdorf et al., 2006; Anton-Erxleben et al., 2009; Klein et al., 2014; Kay et al., 2015; Vo et al., 2017; van Es et al., 2018*). These changes in neural representation are thought to contribute to behavioral enhancement, but because both gain and changes in spatial properties co-occur in biological systems, it is not possible to disentangle them. Computational models of the visual system allow us to design experiments to independently examine the effects of such changes (*Lindsay and Miller, 2018; Eckstein et al., 2000*).

³³ Shrinkage and shift of receptive fields toward attended targets has been observed in both single unit (*Womelsdorf et al., 2006; Anton-Erxleben et al., 2009*) and population (*Klein et al., 2014; Vo et al., 2017; Fischer and Whitney, 2009; van Es et al., 2018*) activity, and has been suggested to lead to behavioral enhancement through a variety of possible mechanisms (*Anton-Erxleben and Carrasco, 2013*). For example, receptive field changes might magnify the cortical representation of attended regions (*Moran and Desimone, 1985*), select for relevant information (*Anton-Erxleben*

39 *et al., 2009; Sprague and Serences, 2013*), reduce uncertainty about spatial position (*Vo et al., 2017*),
40 increase spatial discriminability (*Kay et al., 2015; Fischer and Whitney, 2009*), or change estimates
41 of perceptual size (*Anton-Erxleben et al., 2007*). Compression of visual space is also observed just
42 prior to saccades and thought to shift receptive fields towards the saccade location (*Zirnsak et al.,*
43 *2014; Colby and Goldberg, 1999; Merriam et al., 2007*) and maintain a stable representation of
44 visual space (*Kusunoki and Goldberg, 2003; Tolias et al., 2001; Ross et al., 1997; Duhamel et al.,*
45 *1992*).

46 Shrinkage and shift of receptive fields has also been hypothesized to occur as a side effect of
47 increasing gain of neural responses (*Klein et al., 2014; Compte and Wang, 2006*), thus raising the
48 question of which of these physiological effects could be responsible for enhanced perception.
49 When gain is asymmetric across a receptive field, the overall effect will be to shift the receptive
50 field location towards the side with the largest gain. Similarly, asymmetric gain can be expected
51 to change spatial tuning properties such as the size and structure of the receptive field. These
52 concomitant changes of receptive field size, location, and structure could improve perceptual per-
53 formance through the mechanisms described above, or could be an epiphenomenological conse-
54 quence of increasing gain. Increasing gain by itself has also been hypothesized to be a mechanism
55 for improving perceptual performance, because response gain can increase the signal-to-noise
56 ratio and make responses to different stimuli more discriminable (*McAdams and Maunsell, 1999;*
57 *Cohen and Newsome, 2008*). Moreover, larger responses for attended stimuli can act to select relevant
58 information when read-out through winner-take-all mechanisms (*Lee et al., 1999; Pelli, 1985;*
59 *Pestilli et al., 2011; Palmer et al., 2000; Hara et al., 2014*).

60 We took a modeling approach to ask what effects gain changes incur on spatial receptive field
61 structure when introduced at the earliest stage of visual processing and to ask which effects would
62 improve behavioral performance. We modified a convolutional neural network (CNN) trained on
63 ImageNet categorization to test various hypotheses by implementing them as elements of the
64 model architecture. CNN architectures can be designed to closely mimic the primate visual hier-
65 archy (*Yamins et al., 2014; Kubilius et al., 2018*). Training “units” in these networks to categorize
66 images leads to visual filters that show a striking qualitative resemblance to the filters observed in
67 early visual cortex (*Krizhevsky et al., 2012*) and the pattern of activity of these units when presented
68 with natural images is sufficient to capture a large portion of the variance in neural activity in the
69 retina (*McIntosh et al., 2016*), in early visual cortex (*Cadena et al., 2019*), and in later areas (*Güçlü*
70 *and van Gerven, 2015; Cichy et al., 2016; Eickenberg et al., 2017; Khaligh-Razavi and Kriegeskorte,*
71 *2014; Yamins et al., 2014*). Cortical responses and neural network activity also share a correlation
72 structure across natural image categories (*Storrs et al., 2020*). These properties of CNNs make
73 them a useful tool which we can use to indirectly study visual cortex, probing activity and behavior
74 in ways that are impractical in humans and non-human primates (*Lindsay and Miller, 2018*).

75 Using simulations based on a CNN observer model we found that gain changes introduced at
76 the earliest stage in visual processing improved task performance with a magnitude comparable
77 to that measured in human subjects. While these gain changes also induced changes in receptive
78 field location, size, and spatial structure similar to that reported in physiological measurements,
79 these changes were neither necessary nor sufficient for improving model task performance. More
80 specifically, we designed a simple cued object-detection task and measured improved human per-
81 formance on trials with focal attention. Using CORnet-Z (*Kubilius et al., 2018*), a CNN whose archi-
82 tecture was designed to maximize similarity with the primate visual stream, we measured a similar
83 improvement in detection performance when a Gaussian gain augmented inputs coming from a
84 “cued” location. We found that the network mirrored the physiology of human and non-human
85 primates: units shifted their center-of-mass toward the locus of attention and shrank in size, all in
86 a gain-dependent manner. We isolated each of these physiological changes to determine which, if
87 any, could account for the benefits to performance. A model with only gain reproduced the benefits
88 of cued attention while models with only receptive field shifts, shrinkage, or only changes in recep-
89 tive field structure were unable to provide any benefit to task performance. These results held for

90 both an object detection task and a category discrimination task. Gain applied or removed at the
91 last stage of processing in the CNN observer model demonstrated that gain was both necessary
92 and sufficient to account for the benefits in task performance of the model.

93 **Results**

94 We characterized the ability of human observers to detect objects in a grid of four images, with or
95 without prior information about the object's possible location (Fig. 1). Observers were given a writ-
96 ten category label, e.g. "ferris wheel", and shown five exemplar images of that category (Category
97 intro, 1a). This was followed by a block of 80 trials in which observers tried to detect the presence
98 or absence of the target category among the four images in the grid (Each trial, 1a). Half of the
99 80 trials had focal cues and 50% of the focal (and distributed) trials included a target image. On
100 focal trials a cue indicated with 100% validity the grid quadrant that could contain a target while
101 on distributed trials no information was given as to where an image of the target category could
102 appear. Distractor images were randomly sampled from the nineteen non-target image categories.
103 Stimulus durations were sampled uniformly from 1 (8.3 ms), 2 (16.7), 4 (33.3), 8 (66.7), 16 (133.3),
104 or 32 (267.7) frames (Stimulus, 8.3 ms per frame, 1a). Image grids were masked before and after
105 stimulus presentation by shuffling the pixel locations in the stimulus images, ensuring that the lu-
106 minance during each trial remained constant. Observers had 2 s to make a response and each trial
107 was followed by a 0.25 s inter-trial interval. Observers completed one training block on an unused
108 category prior to data collection.

109 Human observers improved their performance on this detection task when given a focal cue
110 indicating the potential location of a target (Fig. 1b). We quantified human performance by comput-
111 ing sensitivity, d' , as a function of stimulus duration separately for focal and distributed conditions.
112 Across all observers the d' function was best fit as:

$$d'(ms) = \alpha \log(163.6ms + 1) \quad (1)$$

113 Where α scaled the function for the focal condition. At a stimulus duration of 8.3 ms (one frame)
114 observers were near chance performance regardless of cueing condition. On distributed trials
115 observers exceeded threshold performance ($d' = 1$) at a stimulus duration of 155 ms, 95% CI [135,
116 197]. For focal trials, the same threshold was reached with only a 38 ms [32, 43] stimulus duration,
117 demonstrating a substantial performance benefit of the focal cue. We found that d' in the focal
118 condition was higher than in the distributed condition, average increase across observers $\alpha = 1.67 \times$
119 [1.57, 1.74].

120 Using a drift diffusion model we found that the majority of this performance benefit came from
121 improved perceptual sensitivity, rather than speed-accuracy trade off. We assessed this by fitting
122 a drift diffusion model to the reaction time and choice data (*Wagenmakers et al., 2007*). Drift
123 diffusion models assume that responses are generated by a diffusion process in which evidence
124 accumulates over time toward a bound. We used the equations in *Wagenmakers et al. (2007)* to
125 transform each observer's percent correct, mean reaction time, and reaction time variance for the
126 twenty categories and two focal conditions into drift rate, bound separation, and non-decision time.
127 The drift rate parameter is designed to isolate the effect of external input, the non-decision time
128 reflects an internal delay before stimulus processing begins, and the bound separation is a proxy
129 for how conservative observers are. Comparing the drift rate parameter we observed a similar
130 effect to what was described above for d' : the average drift rate across observers in the focal con-
131 dition was $1.61 \times$, 95% CI [1.39, 1.77] the drift rate in the distributed condition. This suggests that
132 the majority of the performance gain observed in the d' parameter came from increased stimulus
133 information. We did find that the other parameters of the drift diffusion model were also sensitive
134 to duration and condition, but in opposite directions. We found larger bound separation at longer
135 stimulus durations and on focal trials (focal bound-separation 1.57× distributed [1.37, 1.75]), con-
136 sistent with observers being more conservative on trials where more information was available.

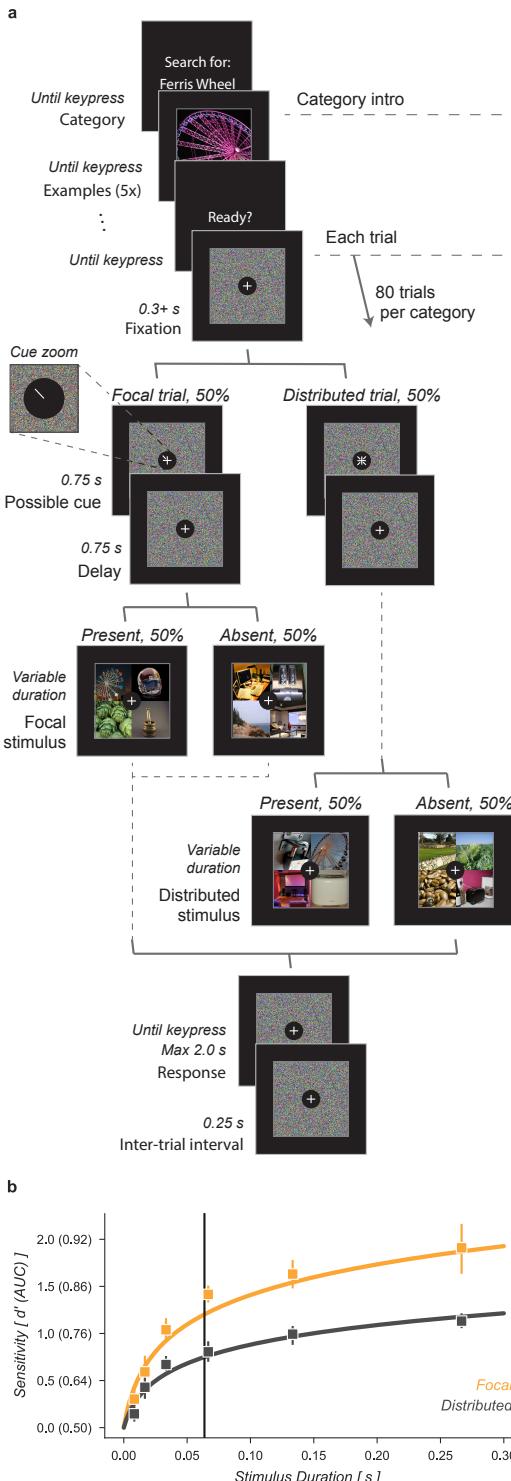


Figure 1. Cued object detection task. (a) Observers were asked to perform object detection with or without a spatial cue. At the start of a block, observers were shown five examples of the target category. This was followed by 80 trials: 40 with a spatial cue indicating the possible target quadrant and 40 with no prior information. Stimulus presentation was pre and post-masked. The stimuli consisted of a composite grid of four individual object exemplars. The target category was present in 50% of trials and always in the cued location on focal trials. Human observers used a keyboard to make a fast button response to indicate the target presence before moving on to the next trial. (b) Human observers showed a substantial improvement in performance when given a focal cue indicating the quadrant at which the target might appear. Vertical line at 64 ms indicates the duration at which the best-fit d' curve for the Distributed condition matched the CNN observer model performance without gain. Markers indicate the median and error bars the 95% confidence intervals.

137 But this increase in cautiousness was offset by a shorter non-decision time on focal trials (0.26 s)
138 compared to distributed (0.38, [0.34, 0.41]).

139 Having shown that a spatial cue provides human observers with increased stimulus information
140 in this task, we next sought to show that a neural network model of the human visual stream could
141 replicate this behavior under similar conditions. We used a convolutional neural network (CNN)
142 model, CORnet-Z (*Kubilius et al., 2018*), a neural network designed to mimic primate V1, V2, V4, and
143 IT and optimized to perform object recognition for images at a similar scale to our task. CORnet-
144 Z is a four layer CNN with repeated convolutional, rectified linear units (ReLU), and pooling (Fig.
145 2d). We used pretrained weights which were optimized for object categorization on ImageNet
146 (*Deng et al., 2009*). To perform our category detection task, we added a fully-connected output
147 layer for each category and trained the weights of that layer to predict the presence of the twenty
148 object categories selected for this study, thus creating a neural network observer model, i.e. a
149 model designed to idealize the computations human observers perform in the 4-quadrant object
150 detection task. We applied the observer model to a task analogous to the one human observers
151 performed (Fig. 2c). The prediction layers added to the end of the model provided independent
152 readouts for the presence or absence of the different target categories (Linear classifier, Fig. 2c).
153 These output layers were trained on a held out set of full-size images from each category. On a
154 separate held out validation set of 100 images, the trained prediction layers achieved a median
155 AUC of 0.90, range [0.77, 0.96].

156 To examine the computational mechanisms that could underlie the performance benefit of
157 focal cues we added a multiplicative Gaussian gain centered at the location of the cued image (Fig.
158 2b, Gaussian width 56 px). We applied this gain at the first layer of the model and tested various
159 strengths of gain.

160 To align the human and model performance we took the performance of the model in the
161 distributed condition (Distributed, Fig. 2a) and found the stimulus duration at which observers
162 in the distributed condition of the human data matched this performance level (64 ms, Fig. 1b).
163 We then scaled up the amplitude of the Gaussian gain incrementally and found that we could
164 mimic the performance enhancement of human spatial attention by setting the maximum of the
165 Gaussian gain field to approximately 4x. The model with this level of gain had a median AUC across
166 categories of 0.80, 95% CI [0.77, 0.82] compared to 0.71 [0.67, 0.72] without gain and a median AUC
167 improvement of 0.09 [0.08, 0.12] within each category.

168 The gain strengths necessary to induce an increase in task performance in the neural network
169 observer model were relatively large compared to the gain due to directed attention observed in
170 measurements of single unit (*Luck et al., 1997; Treue and Trujillo, 1999*) and population (*Birman
171 and Gardner, 2019*) activity. We attribute this difference to the lack of any non-linear "winner-take-
172 all" type of activation in the CNN. In the primate visual system, it is thought that non-linearities
173 such as exponentiation and normalization can accentuate response differences (*Reynolds and
174 Heeger, 2009; Carandini and Heeger, 2012*) and act as a selection mechanism for sensory signals
175 (*Pestilli et al., 2011*). We tested whether similar non-linear mechanisms would allow for smaller
176 gain strengths to be amplified to the range needed by our model. This was tested by raising the
177 activations of units to an exponent before re-normalizing the activation of all units at the output
178 of each layer (see Methods for details). This has the effect of amplifying active units and further
179 suppressing inactive ones. Using this approach we found that a relatively small gain of 1.1x com-
180 bined with an exponent of 3.8 led to a significantly larger effective gain of 1.37x after just one layer
181 (Fig. 3j). This form of non-linearity is consistent with the finding that static output non-linearities
182 in single units range from about 2 to 4 (*Gardner et al., 1999; Albrecht and Hamilton, 1982; Sclar
183 et al., 1990; Heeger, 1992*) and suggests a plausible physiological mechanism by which the larger
184 gains predicted by our model could be implemented. Repeated use of exponentiation and nor-
185 malization in successive layers of the visual system could produce an even larger effective gain. To
186 avoid training a new convolutional neural network (CNN) and possibly violate the close relationship
187 between the primate visual system and the CNN we studied, we continued our analysis without

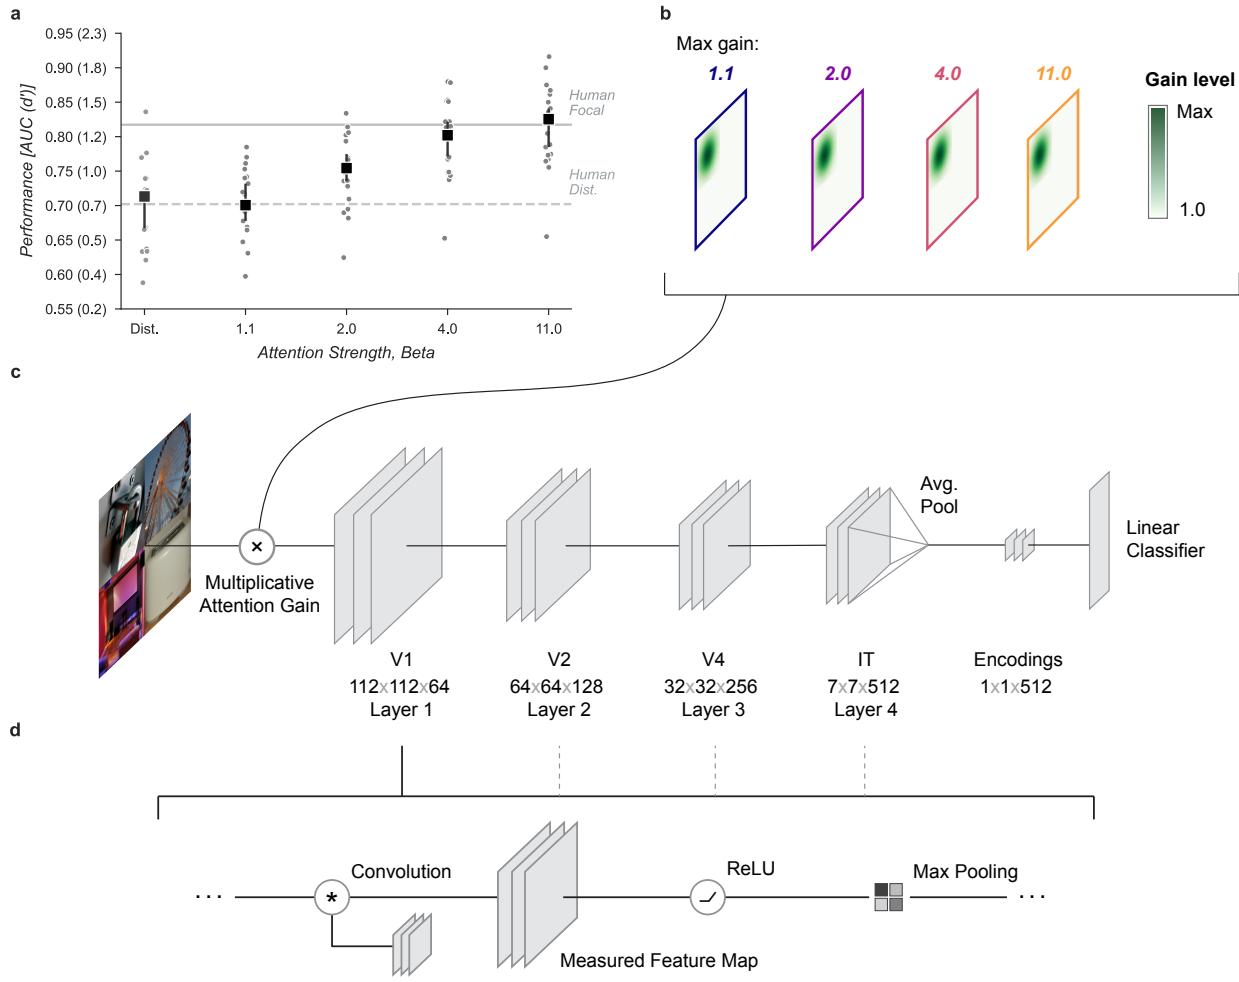


Figure 2. Neural network observer model. (a) Using a Gaussian gain the neural network observer was able to replicate the benefit of spatial attention for human observers. Human performance is shown at a stimulus duration of 64 ms which provided the closest match to the convolutional neural network (CNN) performance without gain. Black markers indicate the median by category and error bars the 95% confidence intervals. (b) The Gaussian gain was implemented by varying the maximum strength of a multiplicative gain map applied to the “cued” quadrant. (c) The gain was applied prior to the first layer of the CNN. The neural network observer model consisted of a four layer CNN with linear classifiers applied to the output layer. Individual classifiers were trained on examples of each object category. (d) Each of the four convolutional layers consisted of a convolution operation, a rectified linear unit, and max pooling. Unit activations were measured at the output of each layer.

188 introducing an exponentiation and normalization step.

189 The Gaussian gain could have its effect on the neural network observer model's performance
190 by increasing the activation strength of units with receptive fields near the locus of attention. These
191 changes in activation strength might directly modify behavior, or work indirectly through mecha-
192 nisms such as changes in receptive field size, location, or spatial tuning. We observed all of these
193 effects in our model (Fig. 3). To measure receptive fields we computed the derivative of each unit
194 with respect to the input image and then fit these with a 2D Gaussian (see Methods for details). We
195 found that the gain caused receptive fields to shift and shrink toward the locus of attention (Fig.
196 3a,b). The information provided by individual units in the model also changed, increasing for units
197 on the border of the cued quadrant (Fig. 3c). The receptive field shift and shrinkage were mag-
198 nified in deeper layers of the model (Fig. 3d,e) consistent with physiological observations (*Klein*
199 *et al.*, 2014). The gain in activation strength propagated through the network without modification
200 (Fig. 3f). To measure the effective gain experienced by the layer four units (Fig. 3i) we computed
201 the ratio of the standard deviations of unit activations at the output of each layer (Fig. 2d) with
202 and without gain applied. All three observed effects: receptive field shift, shrinkage and ex-
203 pansion, and effective gain were directly related to the gain strength at the input layer (Fig. 3g-i). All
204 of these changes have been proposed as mechanisms that could account for the behavioral bene-
205 fits of attention (*Anton-Erxleben and Carrasco*, 2013; *Moran and Desimone*, 1985; *Anton-Erxleben*
206 *et al.*, 2009; *Sprague and Serences*, 2013; *Vo et al.*, 2017; *Kay et al.*, 2015; *Fischer and Whitney*,
207 2009; *Anton-Erxleben et al.*, 2007). We designed models to try to isolate these effects with the goal
208 of testing their independent contributions to task performance.

209 We next sought to test whether receptive field shifts alone could account for the behavioral
210 benefits of the neural network observer model. To do this, we built a model variant that could
211 shift receptive fields without introducing gain. To develop an intuition for how this could affect
212 perceptual reports, consider a CNN with just four units in a 2×2 grid with each unit having its
213 receptive field centered on one image in the composite. When shown a composite grid of four
214 images, a logistic regression using the output of these four units would receive one quarter the
215 information it expects from being trained on full size images. Shifting the receptive fields of the
216 three non-target units to overlap more with the cued image could add additional task-relevant in-
217 formation to the output, much as was observed for units with receptive fields overlapping multiple
218 images in the Gaussian gain attention model (Fig 3c).

219 We designed a variant of our model that could be used to test the hypothesis that receptive field
220 shifts alone are responsible for the behavioral enhancement (Fig. 4). In this model we re-wired the
221 units in the first layer to reproduce the effect of Gaussian gain. The re-wiring was designed so
222 that receptive fields in the fourth layer matched their shift with the Gaussian gain model (Fig. 3g).
223 To mimic those shifts, we changed the connections between the input image pixels and layer one
224 (Fig. 4a). This manipulation worked as designed and changed the receptive field locations and size
225 (Fig. 4b-d) but since no gain was added to the model, the overall responsiveness of units remained
226 constant (Fig. 4e). Because receptive field shifts due to gain are not the result of actual rewiring
227 it is unsurprising that the shift and shrinkage in this model variant are only qualitatively matched
228 to those caused by the original Gaussian gain. Note that the effective gain of individual units in
229 layer four *did* change for individual images, a result of each unit receiving different inputs, but the
230 average change across images was zero.

231 We found that the model with receptive field shifts but no gain had no effect on task perfor-
232 mance, demonstrating that receptive field shifts are not key for the improvement in task perfor-
233 mance observed with Gaussian gain (Fig. 4f). The model imitating shifts from 4x Gaussian gain had
234 a median AUC across categories of 0.71, 95% CI [0.66, 0.73] compared to 0.71 [0.67, 0.72] with no
235 attention and a median change in AUC of -0.01 [-0.02, 0.01] within each category.

236 Another way to understand the possible effect of the Gaussian gain on task performance is to
237 note that the spatial tuning profile of units is "shifted" towards the locus of attention: sensitivity is
238 enhanced closer to the locus of attention, but the receptive field itself has not truly moved in the

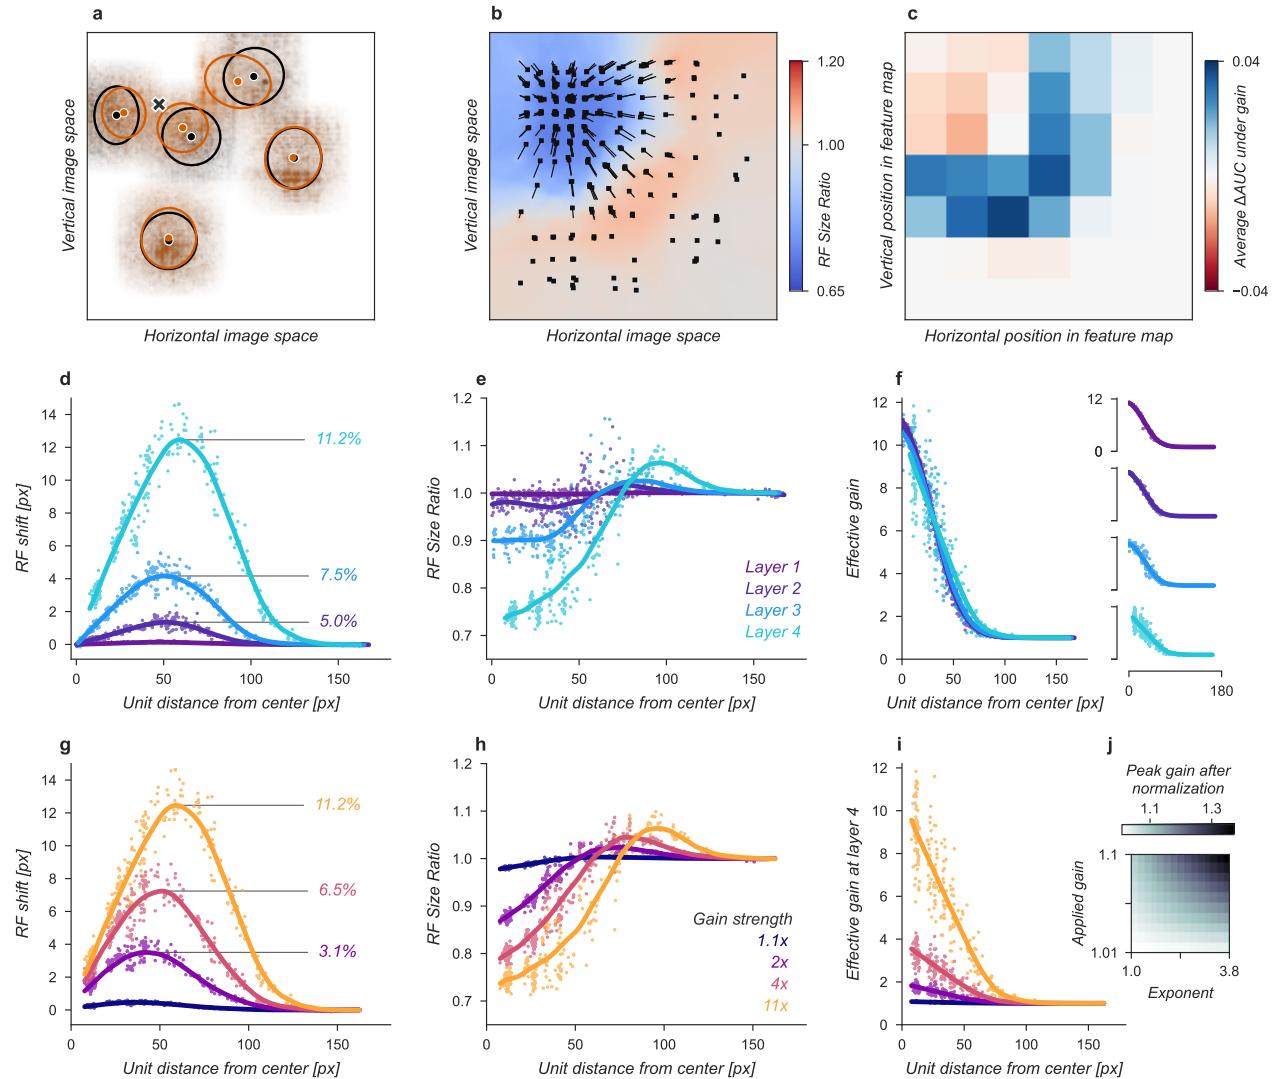


Figure 3. Effects of Gaussian gain on neural network units. (a) The Gaussian gain applied to Layer 1 units caused the measured receptive field (RF) of units in Layer 4 to shift (black ellipse, original; brown ellipse, with gain) toward the locus of attention (black x). (b) A 2D spatial map demonstrates the effects of Gaussian gain in Layer 4: shift of RF center position (black arrows), shrinking RF size near the attended locus (blue colors) and an expansion of size near the gain boundaries (red colors). (c) 7×7 map of the output layer before averaging, showing the change in AUC caused by the addition of Gaussian gain. Each pixel's ΔAUC is computed by projecting the activations at that location for composite grids with target present and absent on the decision axis and then calculating the difference in AUC between a model with and without Gaussian gain. The map demonstrates that units overlapping the borders of the composite grid have the largest change in information content when Gaussian gain is applied. (d,e) Scatter plots demonstrate that each layer magnifies the effect of the gain on RF shift and RF size. The RF shift percentages are the ratio of pixel shift at the peak of the curve relative to the average receptive field size, measured as the full-width at half-maximum. (f) Later layers do not magnify the effective gain (shown for an $11\times$ gain), which stays constant across layers. (g) Gain strength influences the size of RF position shifts, RF size (h), and effective gain (i). (j) Adding an additional non-linear normalizing exponent at the output of each layer allows for much smaller gains to be magnified across layers. Markers in all panels indicate individual sampled units from the model. Lines show the LOESS fit for visualization.

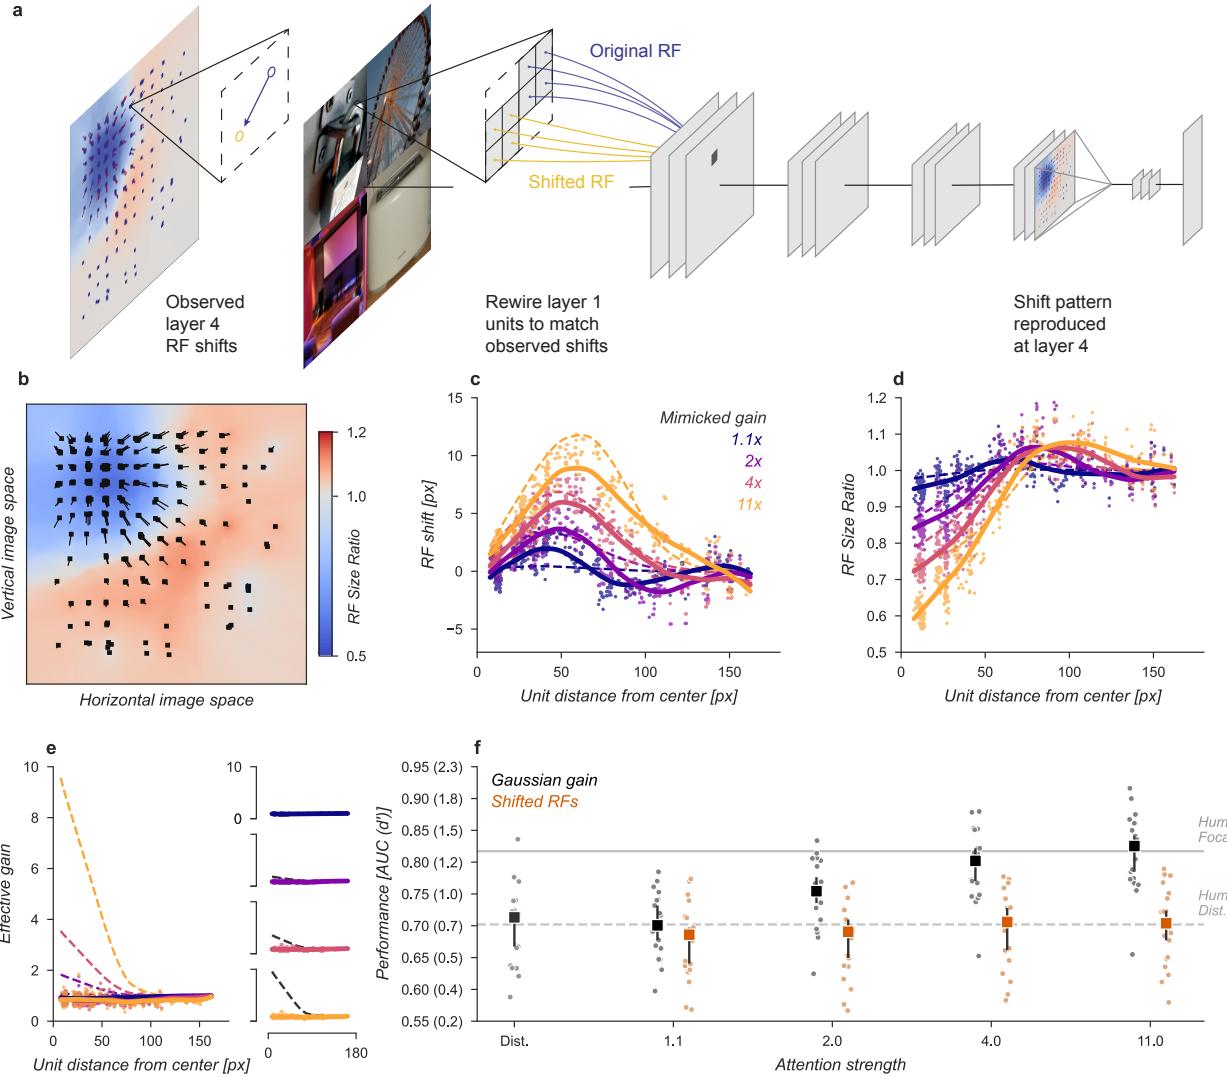


Figure 4. Receptive field shift model. (a) To mimic the effects of the Gaussian gain on receptive field position without inducing gain in the model we re-assigned the inputs to units in Layer 1. This re-assignment was performed so that the pattern of receptive field shift in Layer 4 would match what was observed when the Gaussian gain was applied. (b) The observed pattern of receptive field shifts and shrinkage is shown for a sample of units in layer 4, qualitatively matching the effects of the Gaussian gain. (c) RF shift is shown for sampled units (markers) and the LOESS fit (solid lines) compared to the effect in the Gaussian gain model (dotted lines). (d) Conventions as in c for the RF size change. (e) Conventions as in c,d for the effective gain of units. (f) The behavioral effect of shifting receptive fields is shown to be null on average across categories when compared to the effect of Gaussian gain. Large markers indicate the median performance, small markers the individual categories, and error bars the 95% confidence intervals.

239 manner studied by the previous model. If different parts of a receptive field receive asymmetric
240 gain, as expected for Gaussian gain, then the local structure of the receptive field has been changed
241 (Fig. 5a). We designed another model variant to test the hypothesis that these local changes in re-
242 ceptive field structure might be sufficient to explain changes in task performance without inducing
243 receptive field shifts or gain. To implement this model at layer L , we examined the effect of the
244 Gaussian gain on each unit (green differential gain, Fig. 5a). We normalized this differential gain
245 within each unit's receptive field to prevent any overall gain effect and re-scaled the unit's kernel
246 accordingly. Overall this manipulation of each unit's kernel preserved a portion of the receptive
247 field shift effect in a gain-dependent manner but guaranteed that there was no effective gain.

248 The receptive field structure model was designed to only change the spatial tuning of individual
249 units without inducing gain, which naturally caused some shifts in the measured receptive field size
250 and location (solid lines and markers, Fig. 5b-d) but these were smaller than the effects observed
251 under Gaussian gain (dashed lines). The normalization prevented the model from introducing any
252 spatial pattern of gain change (Fig. 5e). Note that there were still small changes in overall sensitivity
253 of units in this model, for example, the 4x model had an average gain of 1.08, 95% CI [1.07, 1.09]
254 across all units, which we attribute to the fact that inputs to a unit may exhibit correlations due
255 to spatial structure. These receptive field changes and small gain effects were distinct from those
256 observed under Gaussian gain (Fig. 5c-e).

257 The receptive field structure model, like the shift model, was unable to account for the be-
258 havioral effects of the Gaussian gain. No matter where in the model we changed the receptive
259 field structure, and even when applied at all layers, the average performance across categories
260 remained flat (Fig. 5f). Compared to the median distributed AUC across categories of 0.71 [0.67,
261 0.72], the sensitivity model applied to all layers had a median AUC across categories of 0.69 [0.65,
262 0.72] when imitating gain of 1.1x, 0.70 [0.65, 0.72] for 2x gain, 0.69 [0.65, 0.71] for 4x and 0.66
263 [0.63, 0.69] for 11x. Each of these conditions resulted in a median AUC change within category of
264 -0.02 [-0.03, 0.00], -0.01 [-0.03, 0.00], -0.02 [-0.04, -0.01], and -0.04 [-0.05, -0.03], respectively. When
265 applied to early layers we observed a slight drop in performance, which we attribute to how this
266 model directly alters the kernels in the CNN. These changes break the assumption that the CNN
267 kernels at each layer are consistent with those that were optimized when the model weights were
268 trained.

269 The Gaussian gain also caused units to shrink and expand their receptive fields across the vi-
270 sual field (Fig. 3b). These changes might modify the information content received at the output
271 layer, improving or hurting performance. We designed a model variant to test the hypothesis that
272 shrinkage and expansion of receptive fields, without shift or gain, might be sufficient to explain
273 the behavioral effect (Fig. 6). To implement this model we took the observed change in receptive
274 field size at layer 4 and then re-scaled the connections between layers three and four to mimic the
275 observed effect. Because the kernels were scaled in space this manipulation has no effect on ef-
276 fective gain or receptive field position. Specifically, we approximated the shrinkage of the sampled
277 units using a parameterized equation (Eqn. 5) that provides a shrinkage factor for every unit in
278 the model (Fig. 6a). We then re-wired the connections between layer three and four using linear
279 interpolation to approximate the necessary change in scaling.

280 After re-wiring, units' receptive fields retained the same overall position, but were scaled to qual-
281 itatively match the observed effects under Gaussian gain (Fig. 6b-d). The size changes don't match
282 perfectly with those under Gaussian gain because we enforced symmetry in two ways: first, by pa-
283 rameterizing the shrinkage and expansion we enforced symmetry around the locus of attention,
284 and second, because the observed receptive field changes were often asymmetric but we imple-
285 mented a symmetric linear scaling. These simplifications were necessary to reduce the complexity
286 of implementation. By design, the shrinkage and expansion effects scaled by attention strength
287 (Fig. 6d) and induced neither gain-dependent shifts (Fig. 6c) or effective gain (Fig. 6e).

288 The shrinkage model was unable to account for improved task performance with Gaussian
289 gain. The average performance across categories remained flat (Fig. 6f). Compared to the median

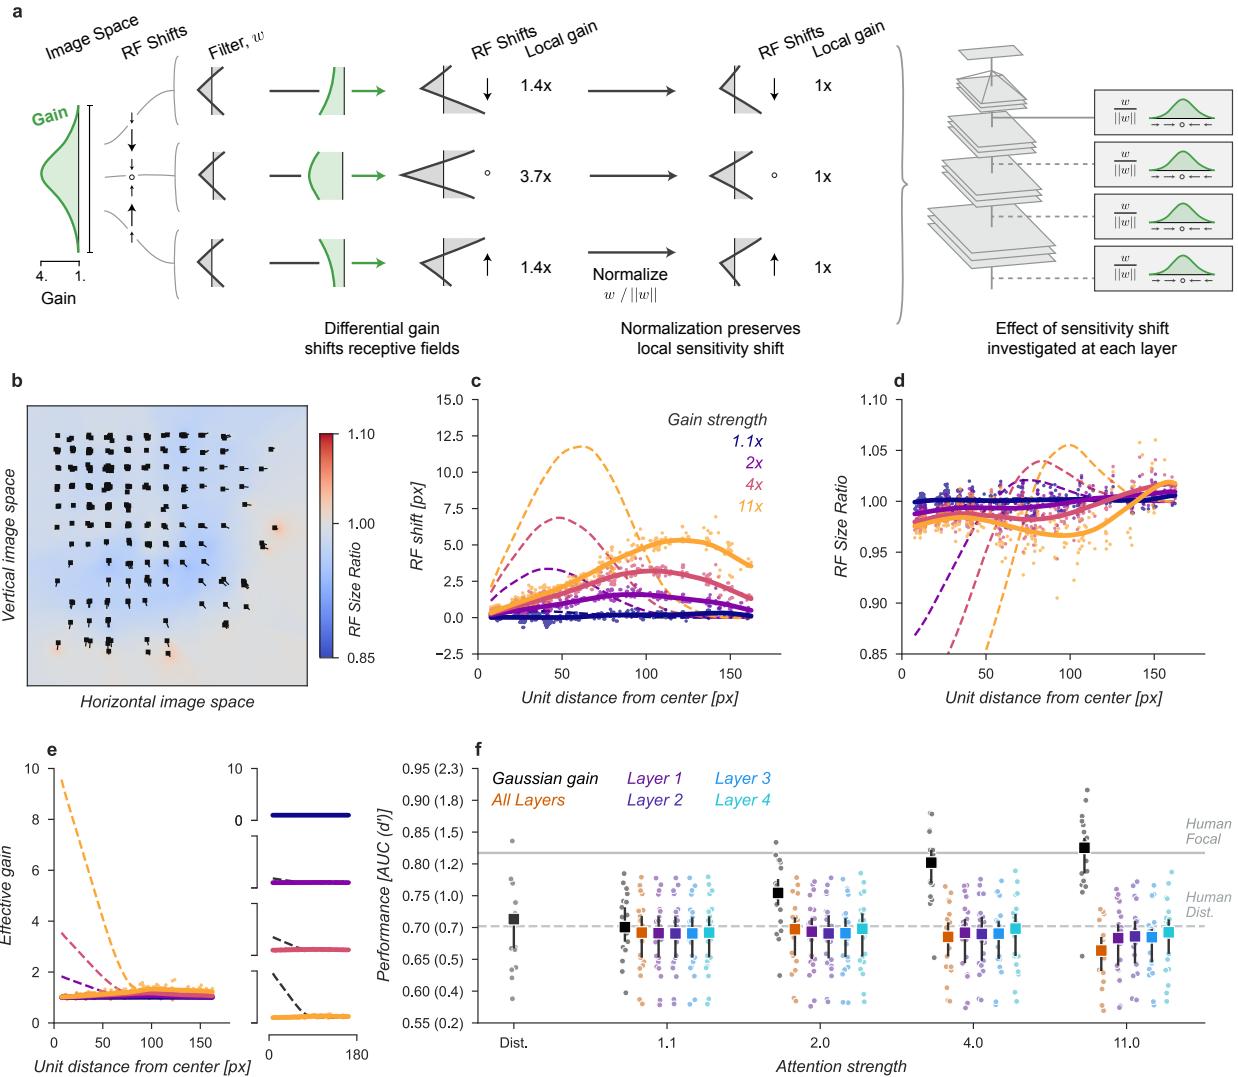


Figure 5. Receptive field structure model. (a) We adjusted the kernels of each convolutional neural network (CNN) unit according to the effect of a Gaussian gain, subtly shifting the the sensitivity within individual units. To avoid inducing a gain change we then normalized each units output such that the sum-of-squares of the weights was held constant, ensuring the local gain at that unit remained at 1x. This model was implemented individually at each layer, replicating the effect of a Gaussian gain of 1.1x to 11x as well as at all layers at once. (b-f) conventions as in Fig. 4.

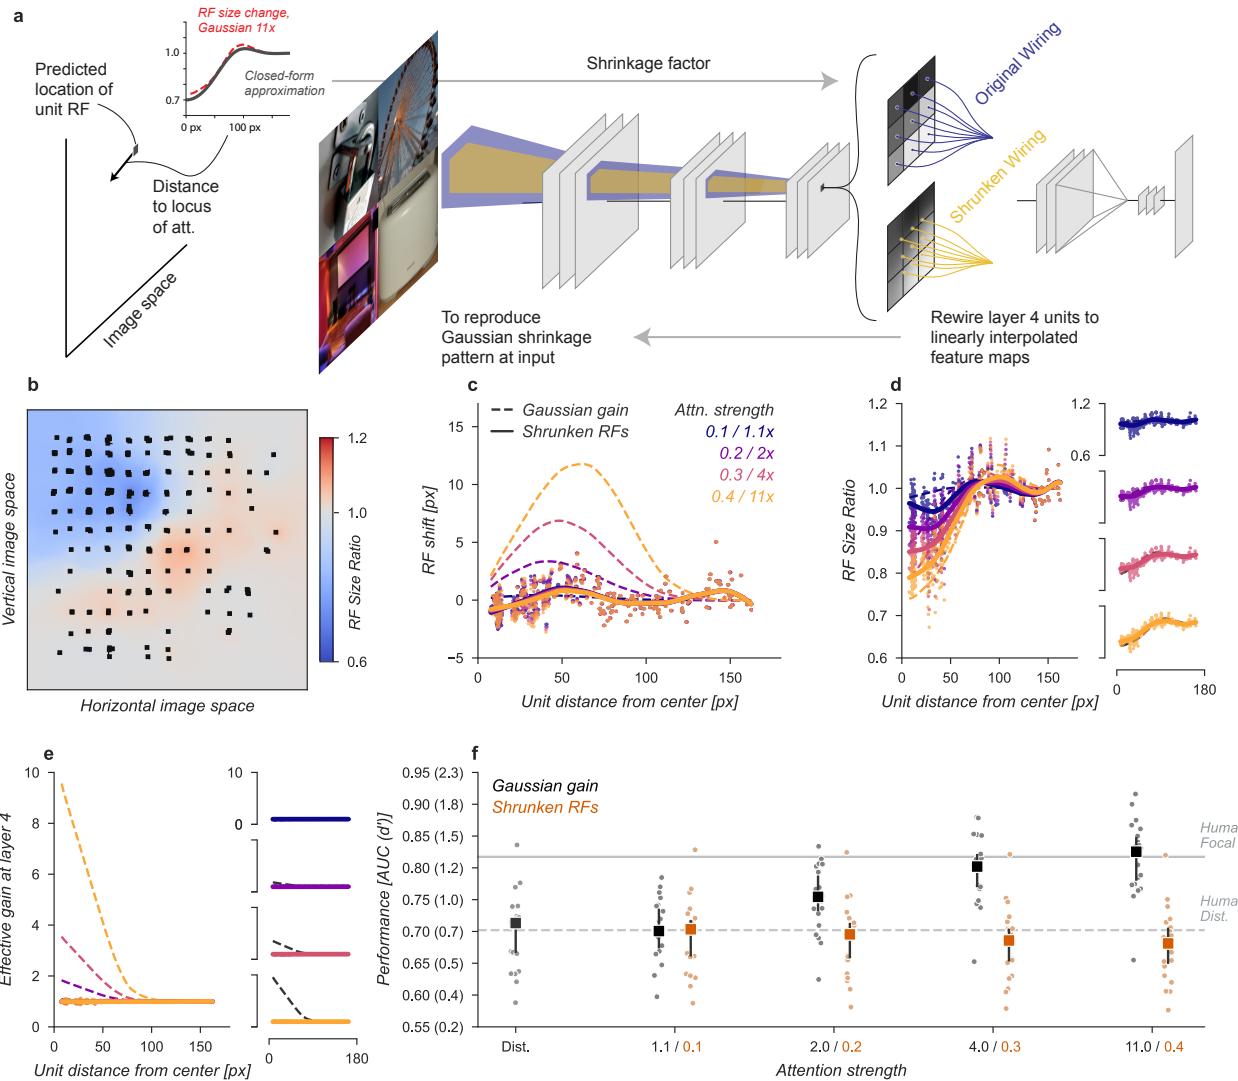


Figure 6. Shrinkage model. (a) To create shrinkage at layer 4 matched with the effects observed under Gaussian gain we re-assigned the connections between layers 3 and 4 according to a parameterized approximation of the shrinkage effect as a function of distance from the locus of attention. This re-scaling of connections changed the size of receptive fields without moving them in space or modifying their gain. (b-f) conventions as in previous figures.

distributed AUC across categories of 0.71 [0.67, 0.72], the shrinkage model applied to all layers had a median AUC across categories of 0.70 [0.66, 0.72] when imitating gain of 1.1 \times , 0.70 [0.66, 0.71] for 2 \times gain, 0.69 [0.65, 0.70] for 4 \times and 0.68 [0.65, 0.70] for 11 \times . Each of these conditions resulted in a median AUC change within category of -0.01 [-0.01, -0.00], -0.01 [-0.01, -0.01], -0.02 [-0.02, -0.01], and -0.02 [-0.03, -0.01], respectively. We again observed drops in performance, which we attribute to how the kernels have been altered.

Having ruled out that receptive field shift, shrinkage, or changes in spatial tuning could account for the improved task performance in our neural network observer, we next designed a model to amplify signals in the cued quadrant without these other effects and found that this model was able to explain the improved task performance observed with cued attention. In the original Gaussian gain model an asymmetry in gain was introduced in the receptive fields of the units, causing size and location changes in the receptive fields. To remove this effect, we flattened the gain within the cued quadrant (Fig. 7a) by setting the gain at each pixel to the average of the Gaussian gain across the entire quadrant. By itself, this change has the unintended consequence that units partly overlapping the cued quadrant will still shift in a gain-dependent manner. To remove this effect, we split the CNN feature maps into the four quadrants and computed these separately with padding and concatenated the results. This forces all units in the model to receive information about only a single quadrant. These manipulations did result in shifts in receptive field location and size for units at the borders (Fig. 7b-d), but by design these were independent of the gain strength.

Using the gain-only model we were able to reproduce the improved task performance of the original Gaussian gain (Fig. 7). The gain-only model induced the same pattern of receptive field shift and size change at all gain strengths (Fig. 7b-d) and a flat effective gain within the cued quadrant (Fig. 7e). We found that increasing the strength of a flat gain was sufficient to capture the full performance improvement of the original model (Fig. 7f). The median AUC across categories of the 4 \times flat gain model was 0.78, 95% CI [0.76, 0.83] compared to 0.80 [0.77, 0.82] for the 4 \times Gaussian gain model. The confidence intervals in flat gain and Gaussian gain performance overlapped at all gain strengths, with a difference of 0.00 [-0.00, 0.02] at 1.1 \times gain, -0.01 [-0.02, 0.00] at 2 \times gain, -0.01 [-0.02, 0.00] at 4 \times gain, and 0.02 [0.00, 0.04] at 11 \times gain.

Having found that the improved task performance could be explained not by receptive field changes, but instead by the change in the overall gain, we asked whether gain propagated through the network was both necessary and sufficient to explain this effect. To test necessity and sufficiency we ran the task images through the Gaussian gain model (first row, Fig. 8a) and measured the effective gain propagated to units in the final layer output ($7 \times 7 \times 512$, before averaging). We averaged these effective gains over features to obtain a propagated gain map (Layer 4 feature map, 7×7 , Fig. 8b). To test the hypothesis that this propagated gain was sufficient to account for improved performance in the task we re-applied it to the output layer of a model with no gain applied to the inputs.

We found that the propagated gain map, when used to multiply the outputs of a model with no Gaussian gain (Multiply by propagated gain, Fig. 8a) was sufficient to induce task performance benefits similar to Gaussian gain applied to the input (Propagated gain vs. Gaussian gain, Fig. 8c). The median AUC across categories using the propagated gain map was 0.79, 95% CI [0.76, 0.84], compared to 0.71 [0.67, 0.72] in the distributed model. There was a small difference between the Gaussian gain and the effect of the propagated gain map -0.02 [-0.03, 0.01], within the 95% confidence interval for no difference. This difference could be attributed to changes in receptive field structure in the Gaussian gain condition, but we attribute it instead to differences between the propagated gain map and the effect of the Gaussian gain. The propagated gain manipulation was constructed from the average effective gain of units across all task stimuli. Because of this, the gain map did not exactly reproduce the effect of gain on an image-by-image basis.

To test the hypothesis that gain was necessary to account for the behavioral effect we divided the final layer activations by the propagated gain map (Divide by propagated gain, Fig. 8a). We found that the behavioral effect of an early gain was mostly reversed by this manipulation (Re-

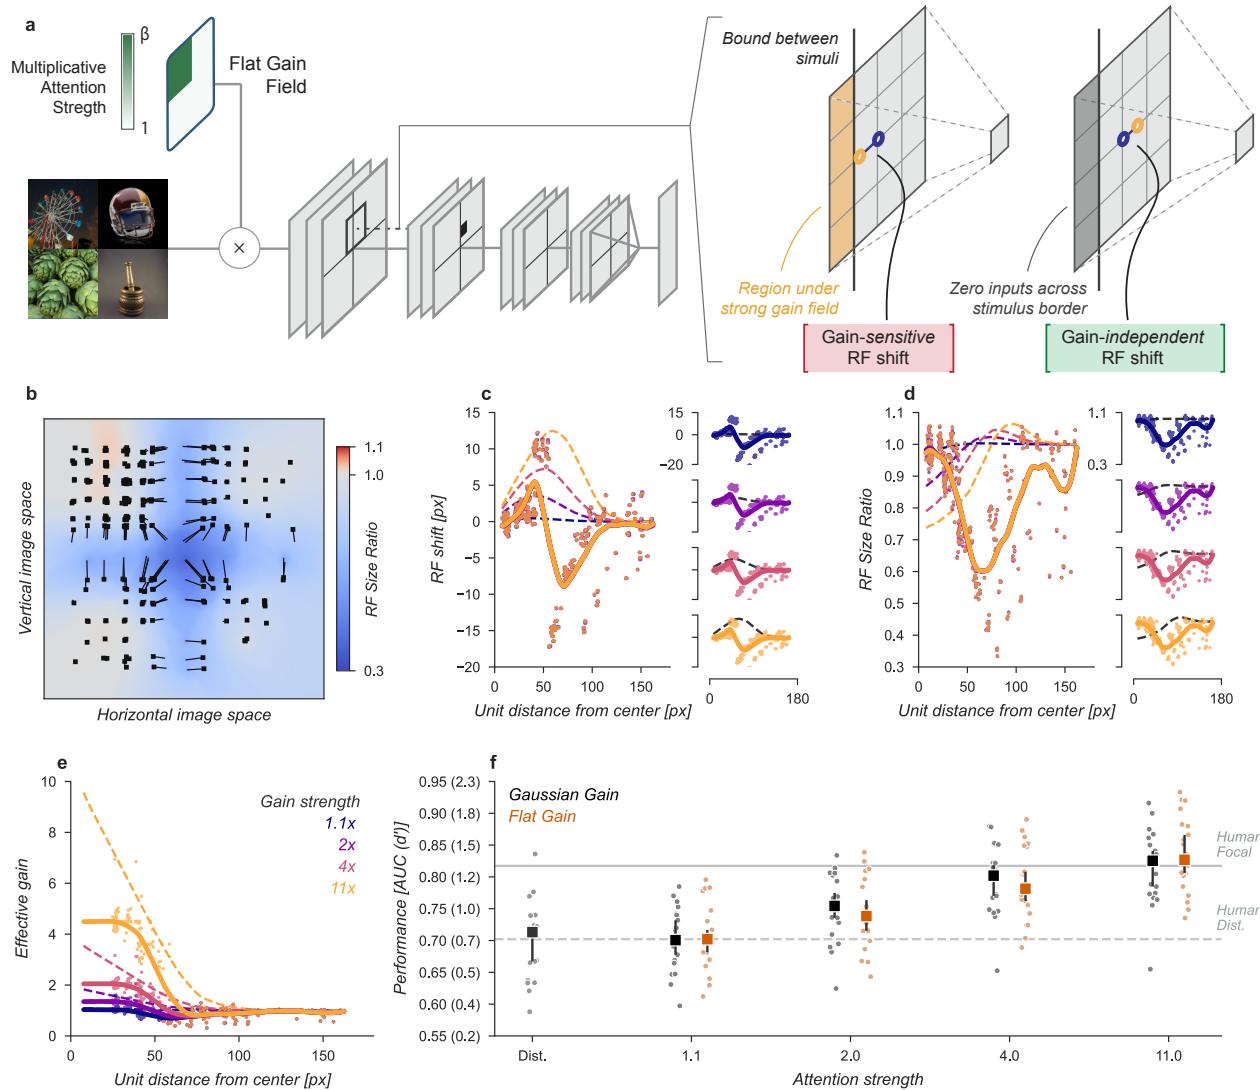


Figure 7. Gain-only model. (a) To create a gain effect without modifying the receptive fields of units we applied a flattened gain field, with the gain set to the average of the original Gaussian gain for each attention strength. The flat gain alone causes units to shift their receptive field at the boundary between the four stimulus quadrants. To modify gain while ensuring shifts were gain-independent we computed the four quadrants separately with zero padding and then concatenated the results. (b-f) conventions as in previous figures.

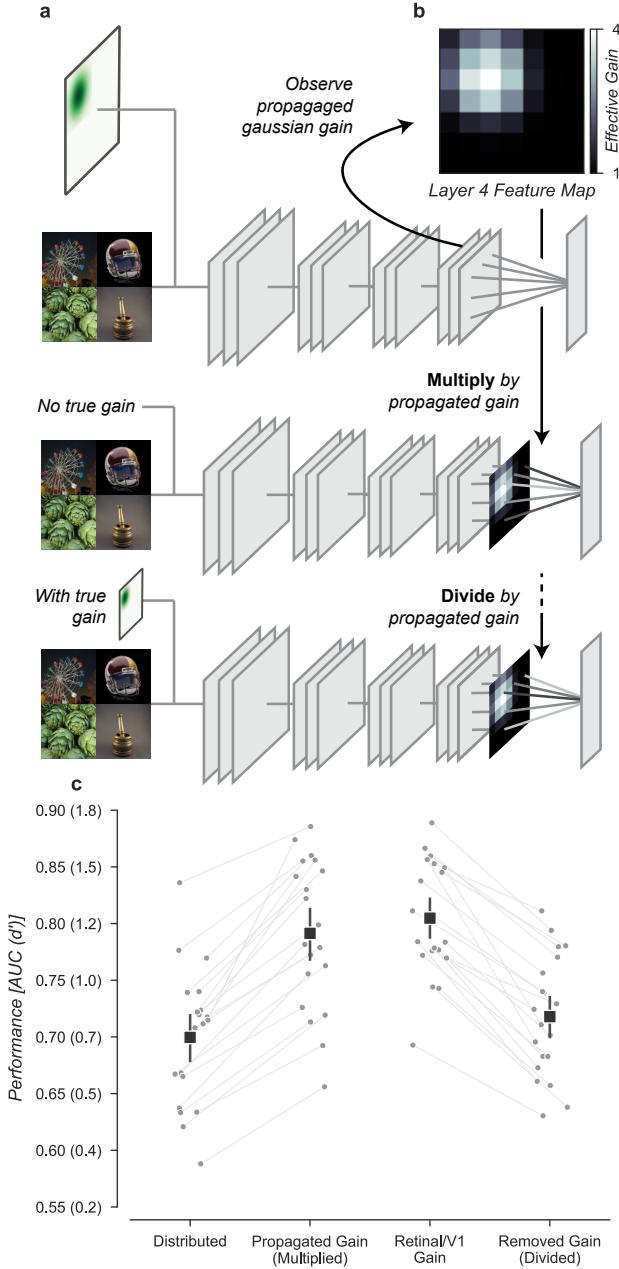


Figure 8. Gain is both necessary and sufficient to explain the improved task performance due to cued attention. (a) To test necessity and sufficiency of gain on performance we propagated the effect of Gaussian gain through the model and measured the effective gain at the output layer. (b) We averaged the effective gain across features to obtain a “propagated gain map”. To test sufficiency we multiplied the output of a model with no true gain by the propagated gain map. To test necessity we divided the output of a model with true gain by the propagated gain map. (c) Multiplying the output by the propagated gain recovered the effect of Gaussian gain, while dividing removed this effect, confirming that gain was both necessary and sufficient to account for the change in task performance. Grey markers show the individual category performance, black markers the median across categories and error bars the 95% confidence intervals.

Figure 8—figure supplement 1. Direct readout from the cued quadrant improves performance alone, with no additional improvement from gain.

Figure 8—figure supplement 2. Gain propagation can account for changes in discrimination task performance due to Gaussian gain.

341 moved gain vs. Distributed, Fig. 8c). The median AUC across categories after dividing out the
342 propagated gain was 0.72, 95% CI [0.68, 0.75], compared to 0.71 [0.67, 0.72] in the distributed con-
343 dition. Dividing by the propagated map did not perfectly reverse the effects of the full Gaussian
344 gain. When compared to the distributed baseline, we found a median within category AUC advan-
345 tage for Gaussian gain with division of 0.02 [0.01, 0.03]. These small residual differences are likely
346 due to the combined effects of changes in spatial receptive field properties that are not reversed
347 by the division of the propagated gain map.

348 Note that in the final readout of our model, we assumed that explicit spatial information was
349 lost, as we averaged activations across the 7×7 convolutional units in the final pooling layer. How-
350 ever, some evidence in ventral temporal cortex suggests that there is spatial information available
351 (*Schwarzlose et al., 2008; Carlson et al., 2011*), so we tested a model read out which retained spatial
352 information and found that the necessary and sufficiency results did not show qualitative changes.
353 A model trained to use the full $7 \times 7 \times 512$ output had marginally worse performance than the model
354 built with the average encodings, achieving a median AUC across categories of 0.68 [0.63, 0.71] in
355 the distributed condition and 0.78 [0.75, 0.82] in the focal condition with 4x Gaussian gain at the
356 first layer. We attribute the small difference in task performance compared to the average model
357 to worse generalization: on the validation set the 7×7 model showed a median drop in AUC across
358 categories of -0.02, range [-0.10, 0.00] compared to the average-pooled readouts.

359 We repeated the propagated gain manipulations in the 7×7 readout model to confirm the
360 necessary and sufficiency results would not change when the model retained spatial information
361 in the final readout. Both the necessity and sufficiency tests showed similar results when using the
362 full output: the average increase in AUC when using the propagated gain map was 0.09, 95% CI
363 [0.08, 0.10] for the full output model, compared to 0.09 [0.07, 0.10] for the average pooled model
364 and the average change in AUC (compared to the distributed condition) when dividing out the
365 propagated gain map from a model with Gaussian gain applied was 0.01 [-0.01, 0.02] for the full
366 output model, compared to 0.02 [0.01, 0.03] for the average pooled model.

367 Improvements in task performance with a Gaussian gain could come from changes in signal
368 discriminability, but also could come from the network being better able to suppress irrelevant
369 visual information. That is, increasing the gain could act to strengthen signals from the relevant
370 target and suppress signals from irrelevant locations. To see how much suppressing irrelevant
371 visual information alone could improve task performance, we designed a neural network observer
372 model which explicitly read out from the top-left $4 \times 4 \times 512$ quadrant of the layer 4 output, instead
373 of the average pooled $1 \times 1 \times 512$ output. As expected, the task performance of this model with no
374 additional gain is already elevated (Distributed, Fig. 8 - figure supplement 1), because the readout
375 now implicitly acts as a form of spatial cueing. The performance of the 4×4 readout was still not
376 at ceiling (Δ AUC between training-validation set and distributed images = -0.07, 95% CI [-0.06, -
377 0.09]). Thus, the performance enhancement due to the Gaussian gain appears to act similarly to
378 an explicit manipulation which suppresses irrelevant information.

379 Theoretical considerations would suggest that moving receptive fields into the target quadrant
380 should further improve performance even when the readout is already spatially specific, because
381 these additional receptive fields can add new information (*Kay et al., 2015; Vo et al., 2017*). We
382 found that this wasn't true for the amount of shift induced by the 4x Gaussian gain, chosen to
383 match the magnitude of the human behavioral benefits of spatial attention. To demonstrate this,
384 we applied a 4x Gaussian gain to the 4×4 readout model and found no further increase in perfor-
385 mance beyond what was achieved by shifting the readout (Gaussian gain vs. 4×4 Readout, Fig. 8 -
386 figure supplement 1). Gain applied to the output of the model also provided no additional benefit
387 (Propagated Gain vs. 4×4 Readout, Fig. 8 - figure supplement 1), supporting the interpretation
388 that the gain acts as a selection mechanism with no effect in the absence of irrelevant distractors.

389 Finally, the observer model solved a detection task where both criterion and sensitivity con-
390 tribute to performance and we reported task performance as AUC to avoid confounding these
391 factors. A more explicit test is to use a criterion-free discrimination task to evaluate the effects of

392 gain on task performance. (Fig. 8 - figure supplement 2a). We therefore designed a category dis-
393 crimination task in which the neural network observer model determined which of two composite
394 grids included the target category at a specified location (always top-left). Baseline performance
395 (Baseline, Fig. 8 - figure supplement 2b) was considered as the performance of the model when no
396 information about which location is cued for discrimination was provided. Note that because the
397 discrimination location always included the target and all other locations had equal probability of
398 including the target category, chance performance was greater than 50%. To compute task perfor-
399 mance, the previously trained fully-connected category target readout was compared across the
400 two composite grids and the composite with the larger response chosen as the model's response.
401 We then applied a Gaussian gain at the cued location and found that discrimination task perfor-
402 mance improved in a similar manner to the detection task (Gaussian Gain, Fig. 8 - figure supple-
403 ment 2b). Using the propagated gain manipulation we confirmed that the gain was both necessary
404 and sufficient for improvements in model task performance (Propagated Gain vs. Gaussian Gain
405 and Removed Gain vs. Chance, Fig. ?? - supplement 2b).

406 Discussion

407 Human observers are more accurate when trying to detect or discriminate objects at a cued loca-
408 tion. Our results demonstrate that this behavioral benefit can also be observed in a neural network
409 model of visual cortex when a Gaussian gain is applied over the pixels of a "cued" object. By mod-
410 eling attentional modulation as gain at the earliest stage of the neural network, we were able to
411 observe similar effects on spatial receptive fields to what is seen in human physiology. When using
412 a gain strength set to match model and human performance enhancement during spatial atten-
413 tion, we documented shifts of receptive fields towards the center of the Gaussian gain field, shrink-
414 age of receptive fields, and changes to the spatial structure in units at later stages of the model.
415 These changes in model receptive field properties were similar in magnitude and characteristics to
416 changes in single-unit (*Womelsdorf et al., 2006; Anton-Erxleben et al., 2009*) and population (*Klein
417 et al., 2014; Vo et al., 2017; Fischer and Whitney, 2009; van Es et al., 2018*) receptive fields reported
418 from physiological measurements.

419 To determine which, if any, of these changes to receptive field properties were the source of
420 improved task performance in the model, we built a series of neural network observer models in
421 which we isolated receptive field shifts, shrinkage, and structural changes from the direct effect
422 of gain. To assess these changes in a way that could provide information about the human visual
423 system, we matched the scale of the shifts, shrinkage, and structural changes to the effect size ob-
424 served in the Gaussian gain model with the gain strength best matched to human performance. In
425 the shift-only model we re-wired units to move receptive fields without introducing gain and found
426 that this produced no improvements in task performance. In the shrinkage model we changed the
427 size of units without changing their gain or position, and again found no improvements in task per-
428 formance. In the receptive field structure model we modified the sensitivity profile of individual
429 receptive fields to mimic the effects of gain, without changing their gain, position, or size, but again
430 found no improvements in task performance. It was only by applying a gain while keeping receptive
431 field properties stable that we were able to reproduce the improvements in task performance.

432 Our results suggest that spatial gain implemented by neural populations in visual cortex can
433 be sufficient to induce behavioral effects of attention for both detection and discrimination even
434 without the concomitant changes in downstream receptive field properties. That is, increasing re-
435 sponse magnitudes through gain changes can act to select relevant visual information when cou-
436 pled with a max or soft-max pooling to suppress irrelevant visual information which has a lower
437 response magnitude (*Lee et al., 1999; Pestilli et al., 2011; Hara et al., 2014; Pelli, 1985*). Although
438 increasing gain can have downstream effects on receptive field properties such as changes in po-
439 sition, size and spatial structure, our results suggest that these may be secondary effects and only
440 a consequence of applying gain, rather than the cause of the behavioral improvements as others
441 have suggested (*Anton-Erxleben and Carrasco, 2013; Moran and Desimone, 1985; Anton-Erxleben*

et al., 2009; Sprague and Serences, 2013; Vo et al., 2017; Kay et al., 2015; Fischer and Whitney, 2009; Anton-Erxleben et al., 2007). In addition, we found that gain had no additional impact when the readout was already spatially specific, reinforcing the interpretation that gain and selection of relevant information are intertwined.

We used an image-computable model of the computational steps from sensory input to decision making which allowed us to formally test hypotheses (*Gardner and Merriam, 2021*) about how different attentional mechanisms could impact task performance. In our case, the advantage of this approach is that the model architecture allowed us to examine how gain at the earliest stages of processing causes changes in spatial receptive field properties: any time a gain occurs in an asymmetrical manner across a receptive field, downstream units will experience an apparent shift as well as shrinkage or expansion. We know from the large literature exploring the physiology of attention that receptive field shifts are correlated with spatial attention (*Anton-Erxleben and Carrasco, 2013; Anton-Erxleben et al., 2009, 2007; Vo et al., 2017; Kay et al., 2015; Fischer and Whitney, 2009; Womelsdorf et al., 2006*). Several authors have proposed that enhanced task performance is a result of increasing information capacity by reducing spatial uncertainty about position (*Kay et al., 2015*) or enhancing discriminability (*Vo et al., 2017*). However, if changes in spatial receptive field properties are the consequence of gain changes (*Klein et al., 2014; Compte and Wang, 2006*), then it raises the question of whether these receptive field changes actually help to improve task performance. Our modeling approach allowed us to examine the theoretical impact of each change associated with gain systematically and quantify that these provided no benefit to detection or discrimination task performance. Nevertheless, previous modeling work has demonstrated that shift and shrinkage, in particular, can increase the resolution and redundancy of receptive field coverage (*Kay et al., 2015; Vo et al., 2017; Theiss et al., 2022*). Our results differ along two important dimensions: first, we scaled the magnitude of shift, shrinkage, and sensitivity changes to those induced by Gaussian gain (at a gain strength that was matched to human performance) and these were in general equal to or smaller in magnitude to what was observed in these previous papers. This leaves open the possibility that larger shifts or shrinkage could have a more direct impact on task performance. We also note that our results were based only on simulations for category detection and discrimination tasks. It may be the case that other tasks which depend more directly on spatial coding, such as judgements of visual position, could exhibit larger benefits from shifts and shrinkage of receptive fields (*Kay et al., 2015; Vo et al., 2017*).

Whether our conclusions can generalize to the behavior of attentional gain in biological neural circuits is limited both by how well the neural network observer model approximates the functioning of those neural circuits and by the model's ability to predict behavior. There are several reasons to suggest that the model captures relevant properties of both object recognition and the primate visual system. We chose to analyze a CNN whose architecture was designed to reflect the primate visual system. This has been evaluated by comparing the similarity of CNN unit activity against measurements of single unit activity in the primate visual cortex (*Schrimpf et al., 2018*). After training, the image features that the CNN units become selective for align closely with those that activate single units in visual cortex (*Yamins et al., 2014; Carter et al., 2019*). In addition, the designers of the architecture we used (CORnet), *Kubilius et al. (2018)* optimized for "core object recognition", detecting a dominant object during a viewing duration of natural fixation (100-200 ms) in the central visual field (10 deg). We re-used core object recognition in our human object detection task and projected our composites in a 10 degree square aperture to obtain similar perceptual characteristics. In the analysis of our task we showed that distributed performance was similar for humans and the CNN at a stimulus presentation of 65 ms, confirming that the intended design of CORnet generalized to the new dataset and task that we used.

While CORnet was designed to map individual visual cortex regions onto the different layers of the CNN, it differs from the visual system in that it is a completely feed-forward model. It is well-known that the visual system has recurrence both within and between visual areas (*Felleman and Van Essen, 1991*). Computational modeling has suggested that recurrence can affect how gain and

493 additive offsets change down stream receptive field location and size, in particular enhancing these
494 effects beyond receptive field boundaries (*Compte and Wang, 2006*). These considerations suggest
495 that more realistic models could have even stronger downstream effects on spatial receptive field
496 properties than what we have documented in a purely feed-forward network. In computational
497 models, recurrent connections are often unfolded into feed-forward layers, effectively making a
498 recurrent model a deeper convolutional model (*Nayebi et al., 2018*). Although we didn't test deeper
499 architectures in our analysis, we expect that the general principles we described should hold for
500 models with more layers and therefore also for models with recurrent connections. An intriguing
501 follow-up direction would be to extend the modeling described here to reaction time tasks, where
502 a recurrent architecture allows for modeling of temporal dynamics and where diffusion models
503 have been found to provide a useful parameterization of how bottom-up and top-down signals
504 contribute to sensory responses over time (*Kay and Yeatman, 2017*).

505 CORnet is also missing many intermediate areas of the visual system (notably area V3) (*Wandell
506 and Winawer, 2011*) as well as an explicit gain control mechanism such as divisive normalization
507 (*Carandini and Heeger, 2012*) which might account for the large gain necessary in our model to
508 produce human-like performance enhancements. These differences mean that the exact strength
509 of the gain signal we observed cannot be mapped directly onto physiology. In particular, while we
510 apply gain at the earliest stage of the model, we do not wish to imply that such a large gain is seen
511 with attention in the LGN inputs to V1 (*O'Connor et al., 2002*). Nor do we imply that the gain in
512 various stages of our model should directly map on to the gain observed in physiological measure-
513 ments, which have tended to highlight larger gain changes in intermediate areas like V4 and MT
514 (*Treue and Trujillo, 1999; McAdams and Maunsell, 1999; Moore and Armstrong, 2003*) compared to
515 earlier areas. Instead, in our model, the 4x gain should be interpreted as both an explicit increase
516 in gain as well as an implicit gain due to the effects of normalization (*Reynolds and Heeger, 2009;
517 Carandini and Heeger, 2012*). While normalization models have traditionally been studied in single
518 layer models, our work extends this general approach to consider downstream effects of gain on
519 RF properties. We assessed how these effects might interact in our CNN by demonstrating that
520 a physiologically plausible gain of 1.1x, when accentuated by a divisive gain control mechanism
521 (*Kaiser et al., 2016; Carandini and Heeger, 2012*) and amplified across multiple visual areas (many
522 of which are not included in the CORnet model), could have produced the magnitude of effects
523 necessary for human-level improvements in task performance. This smaller gain is more consis-
524 tent with neural recordings in primates, where gain changes on the order 20-40% (1.2-1.4x) have
525 been measured (*Motter, 1993; Luck et al., 1997; Treue and Trujillo, 1999*).

526 We chose to model gain at the earliest possible point in the system to understand how signal
527 changes propagate through the visual hierarchy and modify receptive field structure. Physiological
528 measurements have found evidence for early gain (*McAdams and Maunsell, 1999; Motter, 1993;
529 Luck et al., 1997*), but it is equally possible that the gain is applied at a late stage close to decision
530 making and signal gains early in visual cortex are a result of backward projections to these areas
531 (*Buffalo et al., 2010; Moore and Armstrong, 2003*). The propagated gain analysis confirms that gain
532 signals with spatial specificity arriving at later stages in processing (*Moore and Armstrong, 2003*)
533 would have similar effects on task performance.

534 To solve the demands of goal-directed visual attention, the human brain has multiple poten-
535 tial mechanisms available. To select for relevant and suppress irrelevant information, sensory re-
536 sponses can be amplified or the tuning of neurons and populations can be shifted to enhance some
537 signals at the cost of others. In addition, these bottom-up sensory changes can be combined with
538 changes in how sensory representations are read out or communicated to downstream regions.
539 In biological systems, these mechanisms are intertwined: as we have shown, changes to early sen-
540 sory signals will have complex effects on the later stages that are used for readout. In an idealized
541 model, the changes that would have the most effect on the readout would be computed by approx-
542 imating their gradients on the decision axis (*Lindsay and Miller, 2018*). However, these gradients
543 are typically computed in models through back-propagation (*Rumelhart et al., 1986*), and it is not

544 known whether or how similar gradients can be computed in biological systems. Here, we have
545 shown using a state-of-the-art model of the visual system that when the neural network observer
546 is matched with human performance during spatial attention some mechanisms can improve task
547 performance, while others cannot. In the limit, shift, shrinkage, and tuning changes in receptive
548 fields must have an impact on sensory representations and therefore on performance. But our re-
549 sults show that in a neural network model and at the scale expected in the primate visual system
550 during goal-directed behavior, these are not sufficient to produce the expected effects of spatial at-
551 tention on task performance. Instead, gain combined with a nonlinear selection mechanism meets
552 the demands imposed by goal-directed visual attention. New techniques that allow for targeting
553 interventions to defined populations of neurons raise the possibility of manipulating gain and top-
554 down signaling to determine the effect on downstream neural response properties and behavior.
555 Such interventions would allow for testing the main prediction of our model: that spatial visual
556 attention relies primarily on changes in gain and not concomitant downstream effects to spatial
557 receptive field properties.

558 Methods and Materials

559 Human observers

560 Seven observers were observers for the experiments (1 female, 6 male, mean age 22 y, range 19-
561 24). All observers except one (who was an author) were naïve to the intent. No observers were
562 excluded during the initial training sessions (see eye-tracking below). Observers completed 1600
563 trials in two 60 minute sessions. Observers wore lenses to correct vision to normal if needed. Proce-
564 dures were approved in advance by the Stanford Institutional Review Board on human participants
565 research and all observers gave prior written informed consent before participating (Protocol IRB-
566 32120).

567 Hardware setup for human observers

568 Visual stimuli were generated using MATLAB (The Mathworks, Inc.) and MGL (*Gardner et al., 2018*).
569 Stimuli were displayed at 60 cm viewing distance on a 22.5 inch VIEWPixx LCD display (resolution
570 of 1900x1200, refresh-rate of 120 Hz) and responses collected via keyboard. Experiments were
571 performed in a darkened room where extraneous sources of light were minimized.

572 Eye-tracking was performed using an infrared video-based eye-tracker at 500 Hz (Eyelink 1000;
573 SR Research). Calibration was performed at the start of each session to get a validation accuracy
574 of less than 1 degree average offset from expected, using a thirteen-point calibration procedure.
575 During training, trials were initiated by fixating the central cross for 0.5 s and canceled on-line
576 when an observer's eye position moved more than 1.5 degree away from the center of the fixation
577 cross for more than 0.3 s. Observers were excluded prior to data collection if we were unable to
578 calibrate the eye tracker to an error of less than 1 degree of visual angle or if their canceled trial
579 rate did not drop to near zero. All observers passed these criteria. During data collection the online
580 cancellation was disabled and trials were excluded if observers made a saccade outside of fixation
581 (> 1.5deg) during the stimulus period.

582 Experimental Design

583 We compared the ability of humans and neural networks to detect objects in a grid of four images
584 covering 10 degrees of visual angle (224 px). Given a grid of images, the observers were asked
585 to identify whether or not a particular target category was present. On half of the trials we gave
586 observers prior information telling them which of the four grid locations could contain the object
587 (100% valid cue). This focal condition was compared with a distributed condition, in which no
588 information was provided about which grid location could contain the target object. For humans,
589 the prior in the focal condition was a spatial cue, a visual pointer to one corner of the grid. For
590 the neural network, the prior for the focal condition was implemented by a mechanistic change in

591 the model architecture, which differed according to the model of attention being tested. For the
592 neural network, note that in the distributed condition the model is analogous to one in which the
593 focal cue is implemented by a Gaussian of infinite width.

594 To verify that our results were not specific to detection, we also examined the ability of a neural
595 network observer model to perform a category discrimination task. To perform the discrimination
596 we compared the classifier outputs from two composite grids. These grids were constructed such
597 that one of the two grids always contained an image of the target category (A) in the top-left location
598 and the other contained an image from the non-target category (B). The remaining distractors
599 images were randomly sampled from the A and B categories with 50% probability. In the focal cue
600 condition the model architecture was modified to implement a model of attention.

601 **Stimuli: object detection task**

602 In the object detection task, the stimuli presented to both humans and the neural network observer
603 model were composed of four base images arranged in a grid (henceforth a "composite grid"). Each
604 base image contained an exemplar of one of 21 ImageNet (*Deng et al., 2009*) categories. Composite
605 grids always contained images from four different categories. The base images were cropped to be
606 square, and resized to 122 × 122 pixels, making each composite grid 224 × 224 pixels. We pulled 929
607 images from each of 21 ImageNet categories: analog clock (renamed to "clock"), artichoke, bakery
608 (renamed to "baked goods"), banana, bathtub, bonsai tree (renamed to "tree"), cabbage butterfly,
609 coffee, computer, Ferris wheel, football helmet, garden spider (renamed to "spider"), greenhouse,
610 home theater, long-horned beetle (renamed to "beetle"), mortar, padlock, paintbrush, seashore,
611 stone wall, and toaster. These base images were usually representative of their category. However,
612 many included other distracting elements (people, text, strong reflections, etc). Two authors (KF
613 and DB) selected 100 base images for each category absent of distracting elements (low-distraction
614 base images) to be used for the human task. From these low-distraction base images we set aside
615 5 to use as exemplars when introducing the category to human participants.

616 To create the human stimulus set we generated composite grids for each of the 20 target cate-
617 gories. Each category required 80 composite grids: 40 including target objects and 40 without. We
618 therefore needed 40 base images from the target category and 280 (3×40+4×40) base images from
619 the non-target categories. We sampled all images from the low-distraction base images. Targets
620 were placed 10 times in each of the four corners.

621 The neural network observer model was trained and tested on an expanded stimulus set. We
622 set aside 50 base images for each category to train the linear classifiers (see Linear Classifiers,
623 below). The approach was otherwise identical to that described above, but 829 composite grids
624 were created with a target and 829 without, and the composites were assembled from the full set of
625 929 base images. Because CNN models are translation invariant we formed all target composites
626 with the target image in the NW corner, to simplify analysis.

627 **Stimuli: category discrimination task**

628 The stimuli in the category discrimination task were also composite grids of four images. How-
629 ever, these composites were constructed to only include images from a target pair of categories
630 (called "A" and "B" and generated from 20 of the 21 ImageNet categories, as displayed in Table
631 1). Pairs of composites were generated, consisting of an "A" stimulus and a "B" stimulus with the
632 corresponding category in the top left target grid position. The other three locations were filled
633 with distractor images sampled pseudorandomly from the A or B category. Target images were
634 not repeated across composites, but did appear in other stimuli as distractors. We generated 900
635 images per category pair, 450 with an A target and 450 with a B target.

636 **Human object detection task**

637 Human observers performed blocks of trials in which they had to report the presence or absence of
638 a specified category in composite grids. At the start of each block we showed the human observers

Pair	Category A	Category B
0	Ferris wheel	analog clock
1	artichoke	bakery
2	banana	bathtub
3	cabbage butterfly	coffee
4	computer	football helmet
5	garden spider	greenhouse
6	home theatre	long-horned beetle
7	mortar	padlock
8	paintbrush	seashore
9	stone wall	toaster

Table 1. Category pairs for the discrimination task.

the words "Search for:" followed by the name of the current target category (Fig. 1a, Category). They were then shown five held-out (i.e. not shown in the task) exemplar base images to gain familiarity with the target category (Fig. 1a, Examples) and advanced through these with a self-paced button click. This was followed by individual trials of the task. At all times a fixation cross (0.5 deg diameter, white) was visible at the center of the screen in front of a black circle (1 deg diameter). This fixation region obscured the center of the composite grid, but made maintaining fixation easier for observers. At the start of each trial the pixels of the current composite grid were scrambled to create a luminance-matched visual mask. This was displayed until an observer maintained fixation for 0.3 s (Fig. 1a, "Fixation"). Once fixation was acquired a cue was shown for 0.75 s, informing the observer about whether the trial was focal (in which case the possible target location was indicated) or distributed (four possible target locations indicated). The focal cue was a 0.25 deg length white line pointing toward the cued corner of the grid. The distributed cue was four 0.25 deg length white lines pointing toward all four corners of the grid. Distributed and focal cues were presented in pseudo-randomized order throughout each block. The cue was followed by a 0.75 s inter-stimulus interval (Fig. 1a, Delay) before the composite grid (10 × 10 deg) was shown for either 1 (8.3 ms), 2 (16.7), 4 (33.3), 8 (66.7), 16 (133.3), or 32 (266.7) video frames (Fig. 1a, Stimulus). The mask then replaced the stimulus and observers were given 2 s to make a response (Fig. 1a, Response), pressing the "1" key for target present or the "2" key for absent. Feedback was given by changing the fixation cross color to green for correct and red for incorrect until the 2 s period elapsed. A 0.25 s inter-trial interval separated trials.

Observers completed one training block (the "tree" category) as practice before data collection began. They then completed each category block (40 focal trials with 20 target present and 20 target absent, and 40 distributed trials with 20 target present and 20 target absent) before moving on to the next category. Block order was pseudo-randomized for each observer. Each block took about five minutes to complete and a break was provided between blocks, as needed. In total the experiment took about two hours, split into two one hour sessions on different days.

665 Neural network observer model

We modeled the ventral visual pathway using CORnet-Z, a convolutional neural network (CNN) proposed by *Kubilius et al. (2018)*. The model consists of four convolutional layers producing feature maps of decreasing spatial resolution (Table 2). The model which we used was pre-trained on ImageNet by the original authors (*Kubilius et al., 2018*). At the last convolutional layer we took the average over the spatial dimensions of each feature map to create the neural network's representation (512-dimensional vector) of the input image.

	Layer Type		Kernel Size	Output Shape	FWHM (px, deg)
Input				$224 \times 224 \times 3$	
V1 Block	conv, stride=2		7x7	112 \times 112 \times 64	11 (0.5)
	ReLU			56 \times 56 \times 64	
	max pool		2x2	56 \times 56 \times 64	
V2 Block	conv		3x3	56 \times 56 \times 128	26.8 (1.21)
	ReLU			28 \times 28 \times 128	
	max pool		2x2	28 \times 28 \times 128	
V4 Block	conv		3x3	28 \times 28 \times 256	55.6 (2.52)
	ReLU			14 \times 14 \times 256	
	max pool		2x2	14 \times 14 \times 256	
IT Block	conv		3x3	14 \times 14 \times 512	111.4 (5.06)
	ReLU			7 \times 7 \times 512	
	max pool		2x2	7 \times 7 \times 512	
Encodings	avg. pool			1 \times 1 \times 512	

Table 2. CORnet-Z structure. Average receptive field (RF) full-width at half-maximum (FWHM) is measured using ellipses fit to the backpropagated gradients of units in a convolutional layer with respect to the input image pixels. 22.4 pixels corresponds to one degree of visual angle (*Kubilius et al., 2018*).

672 **Linear classifiers: object detection task**

673 To allow the neural network observer model to perform an object detection task we trained a set
 674 of linear classifiers on the model output to predict the presence or absence of each of the twenty
 675 target categories. Each of these fully-connected layers received as input the (512-dimensional) feature
 676 output from the CNN and projected these to a scalar output. Weights were fit using logistic regression with L2 regularization, using *scikit-learn* and the *LIBLINEAR* package (*Pedregosa et al., 2011*). We trained the classifiers on a held out set of base images not used to generate the task
 677 grids, using 50 images with the target present and 50 images with the target absent. Training was
 678 evaluated on an independent validation set of 100 images, median AUC 0.90, range [0.77, 0.96].

679 To test model performance in the detection task the observer model was presented with each
 680 of the composite grids in the full image set and the output of the target category's classifier was
 681 computed. We report the model's area under the curve (AUC) as a measure of performance. The
 682 AUC is computed from the area under the curve defined by plotting the false positive rate against
 683 the true positive rate across the full range of possible thresholds (0 to 1). We used the *scikit-learn*
 684 implementation to calculate the AUC for each model. The AUC can be interpreted as the average
 685 probability that target images will be ranked higher by the logistic regression compared to non-
 686 target images, with a value of 0.5 indicating chance performance and a value of 1 indicating perfect
 687 discrimination. In a signal detection framework, an AUC of 0.75 corresponds approximately to a d'
 688 of 1.

691 **Linear classifiers: category discrimination task**

692 To allow the neural network observer model to perform a category discrimination task we repeated
 693 the linear classifier training described above, adding a final step in which the classifier outputs were
 694 compared for two composites. The composite grid producing a higher output was marked as contain-
 695 ing the target category. The classifiers were trained on a held out set of base images not used
 696 to generate the task grids. Because this task is criterion-free, we report the model's accuracy as
 697 a measure of performance. Note that even in the distributed condition the model performance
 698 exceeds chance: this is because in any set of category pair composites the proportion of grid pos-
 699 tions with a target will always be higher when the target image is fixed to one category. On average
 700 across images the proportion of A images in the A targets will be $2.5/4 (1 + 0.5 + 0.5 + 0.5)$.

701 **Spatial attention: Gaussian gain model**

702 To introduce Gaussian gain as a mechanism for spatial attention we multiplied the pixel intensity
 703 of the input image at row r and column c by the magnitude of a 2-dimensional Gaussian, using the
 704 following equation:

$$g_{r_0,c_0,\sigma,\beta}(r,c) = (\beta - 1) \exp\left(-\frac{(r - r_0)^2 + (c - c_0)^2}{2\sigma^2}\right) + 1 \quad (2)$$

705 Where r_0 and c_0 set the row and column location for the center of the gain field and β controls
 706 the strength, i.e. the multiplicative factor at the peak of the Gaussian. The Gaussian was centered
 707 in the cued quadrant and σ was set to 56 pixels (approx 2.5 degrees). We explored four values of
 708 β : 1.1, 2, 4, and 11.

709 **Quantifying the effects of gain on receptive fields and activations**

710 To reduce computational requirements we randomly sampled 300 units per layer (1,200 total units)
 711 for receptive field analysis, with higher density near the attended locus.

712 To determine the location and size of the receptive field of each CNN unit we computed the
 713 derivative of their activation with respect to the pixels in the input image. This derivative was taken
 714 across a batch of 40 task images evenly distributed across categories. The magnitude of deriva-
 715 tives with respect to the red, green and blue channels were summed to create a sensitivity map.
 716 Receptive field location and size were estimated by fitting a 2D Gaussian distribution. The fit was
 717 performed by treating the sensitivity map as an unnormalized probability distribution and choos-
 718 ing the Gaussian with the same mean and covariance matrix as that distribution. Receptive field
 719 location was measured as the mean of the Gaussian fit. We report the full-width at half-maximum
 720 for the receptive field size.

721 To measure the effect of gain on the activation and information content of CNN units we com-
 722 puted the effective gain and the change in AUC across the sampled units. We defined effective gain
 723 as the ratio between the standard deviation of a unit's activity after applying an attention mech-
 724 anism compared to before. We computed the effective gain across all features and all stimuli. To
 725 compute the change in AUC we measured the average change along the prediction layers' deci-
 726 sion axes for each feature map location in layer 4 between the distributed and focal conditions.
 727 More specifically, for each category and each location in the 7×7 feature map, we passed the
 728 512-dimensional encoding vector onto that category's prediction layer just as we did for the 512-
 729 dimensional vector after average pooling. This resulted in two distributions of confidence scores
 730 along the prediction layer's decision axis (one each for target present and absent), the AUC of which
 731 describes the relative amount of information contained in that feature map location pertaining to
 732 discrimination of target present and absent conditions. We then took the difference of AUCs be-
 733 tween focal and distributed conditions averaged across categories in each location.

734 **Nonlinear normalization**

735 In order to test the ability of "winner-take-all" normalization to amplify small gains, we isolated the
 736 first layer of the CNN, and applied nonlinear normalization with exponent ξ . More precisely, if the
 737 output feature map of the first layer had size M rows by N columns by C channels and activations
 738 a_{ijc} , we calculated the normalized outputs:

$$b_{ijc} = \frac{\sum_{k,l,d=1}^{M,N,C} |a_{kld}|}{\sum_{k,l,d=1}^{M,N,C} |a_{kld}|^\xi} a_{ijc}^\xi. \quad (3)$$

739 To measure the resulting amplified gain we applied a small Gaussian gain between 1x and 1.1x
 740 to the input image in the same manner as in the full Gaussian gain model. We then measured the
 741 ratio of average effective gain for units contained entirely within the gain field against the average
 742 effective gain of units entirely outside the attention gain field, for various values of ξ .

743 **Spatial attention: shift-only model**

744 In the Gaussian gain model we applied the gain at layer 1 and observed changes in the model's
 745 detection performance at the output layers. We took a parallel approach here to design a model
 746 that could mimic the receptive field shifts at layer 4 (induced by gain at layer 1) while producing no
 747 systematic effect on response gain. To cause the layer 4 units to observe different parts of the input
 748 image we shifted the connections between pixels in the input image and first layer. We preserved
 749 all other connections, so layer 4 units of the neural network continued to receive information from
 750 the same layer 1 units.

751 To obtain the size of the necessary connection shifts we created a "shift map" in input image
 752 space by measuring the distance and direction that layer 4 units moved when the Gaussian gain
 753 was applied. To make this measurement, we took each input image pixel location (r, c) and calcu-
 754 lated the average receptive field shift of the 20 sampled layer 4 units with the closest receptive
 755 field centers without attention. Because we used a sampling procedure and not the full set of
 756 layer 4 units we weighted the sampled units by their Euclidean distance from the target pixel. To
 757 reduce noise in the shift map we applied a Gaussian blur with $\sigma = 8$ pixels. Using the shift map,
 758 we then re-assigned the connections from the input image to the layer 1 units. The simplest way
 759 to implement this involved swapping the activation of each layer 1 unit with the activation of
 760 the unit at its shifted location. For example, if unit $(75, 75)$ was shifted by $(-10, -10)$ we assigned it
 761 the activation of the unit at $(65, 65)$. To deal with decimal shifts we performed linear interpolation
 762 using neighboring units.

763 **Spatial attention: receptive field structure**

764 In the receptive field structure model we aimed to mimic the spatial tuning changes induced by the
 765 Gaussian gain at a particular layer but without changing the effective gain of units. To accomplish
 766 this, we first computed the true gain propagated to the target layer L by scaling the Gaussian gain
 767 map to the size of layer $L-1$'s feature map. With this change alone the weights of units closer to the
 768 locus of attention are scaled more than the weights farther from the locus, introducing differential
 769 gain. To avoid a change in the overall scale of units' weights, we re-scaled the kernel to match the
 770 L2-norm (sum-of-squares) of the original kernel weights.

771 To summarize, suppose that layer $L-1$'s feature map is t times the size of the input image so
 772 that a unit at row r and column c of the layer $L-1$ feature map has an effective effective gain of
 773 $g_{tr_0,tc_0,\sigma,\beta}(tr, tc)$ under the Gaussian gain model. Then if $w \in \mathbb{R}^N$ is the original weight vector of a
 774 unit in the unraveled convolution at layer L whose input vector $a \in \mathbb{R}^N$ contains the activations of
 775 post-ReLU units of layer $L-1$, and if the row-column positions in the $L-1$ feature map of the unit
 776 described by a_i is (r_i, c_i) , then the replacement weight vector in the sensitivity shift model is given
 777 by the vector $w' \in \mathbb{R}^N$, whose entries are:

$$w'_i = \left(\frac{\sum_{i=1}^N w_i^2}{\sum_{i=1}^N w_i^2 g_{tr_0,tc_0,\sigma,\beta}(tr_i, tc_i)^2} \right)^{1/2} w_i, \quad (4)$$

778 **Spatial attention: Shrinkage model**

779 In the shrinkage model we aimed to mimic the receptive field size changes observed at layer 4
 780 under Gaussian gain, without causing changes in receptive field location or gain. To achieve this, we
 781 assigned a shrinkage factor to each layer 4 unit and rewired its connections to layer 3 accordingly.

782 Shrinkage factors $f_\beta(d)$ were determined by the distance d between the locus of attention in
 783 input image space and the unit's spatial location in the feature map projected back onto the input
 784 image. This distance was converted to a shrinkage factor by a function chosen to model the prop-
 785 erties of the receptive field size change pattern observed under 11xx Gaussian gain at layer 4 (Fig.
 786 3e), namely

$$f_\beta(d) = 1 - \beta \exp\left(-2.44 \frac{d^2}{112^2}\right) \cos\left(2.89 \frac{d^2}{112^2}\right) \quad (5)$$

787 where β determines the overall strength of the effect, and ranged from 0.1 to 0.4 in our analyses.
 788 A shrinkage factor of 0 indicated no change in receptive field size, while a shrinkage factor of 1
 789 indicated shrinkage to zero radius.

790 Given a shrinkage factor, we re-weighted the connections of each layer 4 unit to produce an
 791 approximate shrunken convolutional kernel for that unit. The linearity of convolution provides an
 792 equivalence between re-weighting connections from layer 3 to layer 4 and replacing those connec-
 793 tions with new ones to units in a virtual continuous layer 3 feature map formed by linear interpo-
 794 lation between activations in the true layer 3 feature map. We therefore were able to calculate the
 795 new weights for each layer 4 unit based on a length-9 array of floating-point locations on the layer
 796 3 feature map (all CORnet-Z kernels are 3×3). Given the original wiring locations $x_i, y_i, i = 1, \dots, 9$
 797 for a unit with distance d , the new location corresponding to input i was chosen to be

$$x'_i = f(d)x_i - (1 - f(d))\left(\frac{1}{9} \sum_{j=1}^9 x_j\right) \quad (6)$$

798 and similarly for y'_i . Using the linearity of convolution, each new virtual input location (x'_i, y'_i) is
 799 equivalent (for a linearly interpolated feature map) to a weighted combination of connections to
 800 the four feature map locations surrounding (x'_i, y'_i) , calculated by rounding x and y coordinates up
 801 or down. The resultant 36 (9×4) connections were then simplified by combining connections from
 802 the same layer 3 unit to yield a re-weighted convolution kernel.

803 **Spatial attention: Gain-only model**

804 We designed a model which could effect gain without receptive field shift by flattening the gain in
 805 the cued quadrant. Receptive field shift occurs when there is a differential gain across the receptive
 806 field of a unit. To get rid of this, you can put a flat gain over the cued quadrant. This naive approach
 807 has the problem that units that overlap two quadrants will still shift and shrink according to the
 808 strength of the gain. To prevent these units from shifting in a manner correlated to the gain we
 809 separated the CNN feature maps into four parts corresponding to the four image quadrants, ran
 810 the model forward with zero padding around each quadrant, and then concatenated the results
 811 back together. This ensured that each unit experienced a flat gain across its inputs and that as gain
 812 increased units near the quadrant boundaries did not experience gain-dependent receptive field
 813 shift or shrinkage.

814 **Necessary and sufficient test**

815 To obtain a propagated gain map in the final layer output we applied the Gaussian gain to the
 816 start of the neural network observer model and measured the average effective gain of the 7×7
 817 layer 4 output units across a representative sample of images. We call this the “propagated gain
 818 map”, since it represents the effect of the input gain on the output layers. We tested necessity by
 819 dividing the network output by the map for a model with gain applied and we tested sufficiency by
 820 multiplying the outputs from a no-gain model.

821 **Readout from target quadrant**

822 To test the behavior of the neural network observer model with spatially-specific readout from the
 823 last convolution layer (Layer 4), we masked the output of that layer to the linear prediction layers
 824 in the object detection task. To apply the mask, we zeroed activations of units outside the top-left
 825 $4 \times 4 \times 512$ of layer 4 (full dimensions $7 \times 7 \times 512$). The same linear prediction layers and stimuli were
 826 used as in the necessary and sufficient test, and the same four conditions were tested: no gain,
 827 early Gaussian gain, and with a propagated gain map applied and divided out at layer 4.

828 **Behavioral analysis**

829 We analyzed the human behavioral data by binning trials according to their duration and comput-
830 ing sensitivity d' from the equation:

$$d' = Z(H) - Z(FA) \quad (7)$$

831 Where Z is the inverse of the cumulative normal distribution and H and FA are the hit and
832 false alarm rate, respectively. We fit a logarithmic function to the d' data using the equation:

$$d'(t) = \alpha * \log(\kappa t + 1) \quad (8)$$

833 Where t is the stimulus duration and α and κ are parameters that control the shape of the
834 logarithmic function.

835 To compare human and model performance we can also convert between d' and the area under
836 the curve (AUC) by the equation:

$$d' = \sqrt{2}Z(AUC) \quad (9)$$

837 **Confidence intervals**

838 All error bars are calculated by bootstrapping the given statistic with $n = 1000$ and reported as the
839 95% confidence interval.

840 **Data and code availability**

841 The images and composite grids used in this study as well as the code necessary to replicate our
842 analyses are available in the Open Science Framework with the identifier 10.17605/OSF.IO/AGHQK.

843 **Acknowledgments**

844 We thank Josh Wilson for help with data collection and Eline Kupers and Maggie Henderson for
845 early discussions.

846 **References**

- 847 Albrecht DG, Hamilton DB. Striate cortex of monkey and cat: contrast response function. *Journal of neuro-*
848 *physiology*. 1982; 48(1):217–237.
- 849 Anton-Erxleben K, Carrasco M. Attentional enhancement of spatial resolution: linking behavioural and neuro-
850 *physiological evidence*. *Nature Reviews Neuroscience*. 2013; 14(3):188–200.
- 851 Anton-Erxleben K, Henrich C, Treue S. Attention changes perceived size of moving visual patterns. *Journal of*
852 *Vision*. 2007; 7(11):5–5.
- 853 Anton-Erxleben K, Stephan VM, Treue S. Attention reshapes center-surround receptive field structure in
854 macaque cortical area MT. *Cerebral cortex*. 2009; 19(10):2466–2478.
- 855 Ben Hamed S, Duhamel JR, Bremmer F, Graf W. Visual receptive field modulation in the lateral intraparietal
856 area during attentive fixation and free gaze. *Cerebral cortex*. 2002; 12(3):234–245.
- 857 Birman D, Gardner JL. A flexible readout mechanism of human sensory representations. *Nature communica-*
858 *tions*. 2019; 10(1):1–13.
- 859 Buffalo EA, Fries P, Landman R, Liang H, Desimone R. A backward progression of attentional effects in the
860 ventral stream. *Proceedings of the National Academy of Sciences*. 2010; 107(1):361–365.
- 861 Cadena SA, Denfield GH, Walker EY, Gatys LA, Tolias AS, Bethge M, Ecker AS. Deep convolutional mod-
862 els improve predictions of macaque V1 responses to natural images. *PLoS computational biology*. 2019;
863 15(4):e1006897.
- 864 Carandini M, Heeger DJ. Normalization as a canonical neural computation. *Nature Reviews Neuroscienc*e.
865 2012; 13(1):51–62.

- 866 Carlson T, Hogendoorn H, Fonteijn H, Verstraten FA. Spatial coding and invariance in object-selective cortex.
867 Cortex. 2011; 47(1):14–22.
- 868 Carrasco M. Visual attention: The past 25 years. Vision research. 2011; 51(13):1484–1525.
- 869 Carter S, Armstrong Z, Schubert L, Johnson I, Olah C. Activation Atlas. Distill. 2019; doi: [10.23915/distill.00015](https://doi.org/10.23915/distill.00015),
870 <https://distill.pub/2019/activation-atlas>.
- 871 Cichy RM, Khosla A, Pantazis D, Torralba A, Oliva A. Comparison of deep neural networks to spatio-temporal
872 cortical dynamics of human visual object recognition reveals hierarchical correspondence. Scientific reports.
873 2016; 6(1):1–13.
- 874 Cohen MR, Newsome WT. Context-dependent changes in functional circuitry in visual area MT. Neuron. 2008;
875 60(1):162–173.
- 876 Colby CL, Goldberg ME. Space and attention in parietal cortex. Annual review of neuroscience. 1999; 22(1):319–
877 349.
- 878 Compte A, Wang XJ. Tuning curve shift by attention modulation in cortical neurons: a computational study of
879 its mechanisms. Cerebral Cortex. 2006; 16(6):761–778.
- 880 Connor CE, Gallant JL, Preddie DC, Van Essen DC. Responses in area V4 depend on the spatial relationship
881 between stimulus and attention. Journal of neurophysiology. 1996; 75(3):1306–1308.
- 882 Deng J, Dong W, Socher R, Li LJ, Li K, Fei-Fei L. ImageNet: A Large-Scale Hierarchical Image Database. In: CVPR09;
883 2009..
- 884 Duhamel JR, Colby CL, Goldberg ME. The updating of the representation of visual space in parietal cortex by
885 intended eye movements. Science. 1992; 255(5040):90–92.
- 886 Eckstein MP, Thomas JP, Palmer J, Shimozaki SS. A signal detection model predicts the effects of set size on
887 visual search accuracy for feature, conjunction, triple conjunction, and disjunction displays. Perception &
888 psychophysics. 2000; 62(3):425–451.
- 889 Eickenberg M, Gramfort A, Varoquaux G, Thirion B. Seeing it all: Convolutional network layers map the function
890 of the human visual system. NeuroImage. 2017; 152:184–194.
- 891 van Es DM, Theeuwes J, Knapen T. Spatial sampling in human visual cortex is modulated by both spatial and
892 feature-based attention. Elife. 2018; 7:e36928.
- 893 Felleman DJ, Van Essen DC. Distributed hierarchical processing in the primate cerebral cortex. Cerebral cortex
894 (New York, NY: 1991). 1991; 1(1):1–47.
- 895 Fischer J, Whitney D. Attention narrows position tuning of population responses in V1. Current biology. 2009;
896 19(16):1356–1361.
- 897 Gardner JL, Anzai A, Ohzawa I, Freeman RD. Linear and nonlinear contributions to orientation tuning of simple
898 cells in the cat's striate cortex. Visual neuroscience. 1999; 16(6):1115–1121.
- 899 Gardner JL, Merriam EP. Population Models, Not Analyses, of Human Neuroscience Measurements. Annual
900 Review of Vision Science. 2021; 7:225–255.
- 901 Gardner JL, Merriam EP, Schluppeck D, Larsson J. MGL: Visual psychophysics stimuli and experimental design
902 package. Zenodo. 2018 Jun; .
- 903 Güçlü U, van Gerven MA. Deep neural networks reveal a gradient in the complexity of neural representations
904 across the ventral stream. Journal of Neuroscience. 2015; 35(27):10005–10014.
- 905 Hara Y, Pestilli F, Gardner JL. Differing effects of attention in single-units and populations are well predicted by
906 heterogeneous tuning and the normalization model of attention. Frontiers in computational neuroscience.
907 2014; 8:12.
- 908 Hawkins HL, Hillyard SA, Luck SJ, Mouloua M, Downing CJ, Woodward DP. Visual attention modulates signal
909 detectability. Journal of Experimental Psychology: Human Perception and Performance. 1990; 16(4):802.
- 910 Heeger DJ. Half-squaring in responses of cat striate cells. Visual neuroscience. 1992; 9(5):427–443.

- 911 **Kaiser D**, Oosterhof NN, Peelen MV. The neural dynamics of attentional selection in natural scenes. *Journal of
912 neuroscience*. 2016; 36(41):10522–10528.
- 913 **Kay KN**, Weiner KS, Grill-Spector K. Attention reduces spatial uncertainty in human ventral temporal cortex.
914 *Current Biology*. 2015; 25(5):595–600.
- 915 **Kay KN**, Yeatman JD. Bottom-up and top-down computations in word-and face-selective cortex. *elife*. 2017;
916 6:e22341.
- 917 **Khaligh-Razavi SM**, Kriegeskorte N. Deep supervised, but not unsupervised, models may explain IT cortical
918 representation. *PLoS computational biology*. 2014; 10(11):e1003915.
- 919 **Klein BP**, Harvey BM, Dumoulin SO. Attraction of position preference by spatial attention throughout human
920 visual cortex. *Neuron*. 2014; 84(1):227–237.
- 921 **Krizhevsky A**, Sutskever I, Hinton GE. Imagenet classification with deep convolutional neural networks. *Advances
922 in neural information processing systems*. 2012; 25:1097–1105.
- 923 **Kubilius J**, Schrimpf M, Nayebi A, Bear D, Yamins DL, DiCarlo JJ. Cornet: Modeling the neural mechanisms of
924 core object recognition. *BioRxiv*. 2018; p. 408385.
- 925 **Kusunoki M**, Goldberg ME. The time course of perisaccadic receptive field shifts in the lateral intraparietal
926 area of the monkey. *Journal of neurophysiology*. 2003; 89(3):1519–1527.
- 927 **Lee DK**, Itti L, Koch C, Braun J. Attention activates winner-take-all competition among visual filters. *Nature
928 neuroscience*. 1999; 2(4):375–381.
- 929 **Lindsay GW**, Miller KD. How biological attention mechanisms improve task performance in a large-scale visual
930 system model. *ELife*. 2018; 7:e38105.
- 931 **Luck SJ**, Chelazzi L, Hillyard SA, Desimone R. Neural mechanisms of spatial selective attention in areas V1, V2,
932 and V4 of macaque visual cortex. *Journal of neurophysiology*. 1997; 77(1):24–42.
- 933 **McAdams CJ**, Maunsell JH. Effects of attention on orientation-tuning functions of single neurons in macaque
934 cortical area V4. *Journal of Neuroscience*. 1999; 19(1):431–441.
- 935 **McIntosh LT**, Maheswaranathan N, Nayebi A, Ganguli S, Baccus SA. Deep learning models of the retinal re-
936 sponse to natural scenes. *Advances in neural information processing systems*. 2016; 29:1369.
- 937 **Merriam EP**, Genovese CR, Colby CL. Remapping in human visual cortex. *Journal of neurophysiology*. 2007;
938 97(2):1738–1755.
- 939 **Moore T**, Armstrong KM. Selective gating of visual signals by microstimulation of frontal cortex. *Nature*. 2003;
940 421(6921):370–373.
- 941 **Moran J**, Desimone R. Selective attention gates visual processing in the extrastriate cortex. *Science*. 1985;
942 229(4715):782–784.
- 943 **Motter BC**. Focal attention produces spatially selective processing in visual cortical areas V1, V2, and V4 in the
944 presence of competing stimuli. *Journal of neurophysiology*. 1993; 70(3):909–919.
- 945 **Nayebi A**, Bear D, Kubilius J, Kar K, Ganguli S, Sussillo D, DiCarlo JJ, Yamins DL. Task-driven convolutional
946 recurrent models of the visual system. *Advances in neural information processing systems*. 2018; 31.
- 947 **O'Connor DH**, Fukui MM, Pinsk MA, Kastner S. Attention modulates responses in the human lateral geniculate
948 nucleus. *Nature neuroscience*. 2002; 5(11):1203–1209.
- 949 **Palmer J**, Verghese P, Pavel M. The psychophysics of visual search. *Vision research*. 2000; 40(10-12):1227–1268.
- 950 **Pedregosa F**, Varoquaux G, Gramfort A, Michel V, Thirion B, Grisel O, Blondel M, Prettenhofer P, Weiss R,
951 Dubourg V, et al. Scikit-learn: Machine learning in Python. *the Journal of machine Learning research*. 2011;
952 12:2825–2830.
- 953 **Pelli DG**. Uncertainty explains many aspects of visual contrast detection and discrimination. *JOSA A*. 1985;
954 2(9):1508–1532.
- 955 **Pestilli F**, Carrasco M, Heeger DJ, Gardner JL. Attentional enhancement via selection and pooling of early sen-
956 sory responses in human visual cortex. *Neuron*. 2011; 72(5):832–846.

- 957 **Posner MI**. Orienting of attention. *Quarterly journal of experimental psychology*. 1980; 32(1):3–25.
- 958 **Reynolds JH**, Heeger DJ. The normalization model of attention. *Neuron*. 2009; 61(2):168–185.
- 959 **Ross J**, Morrone MC, Burr DC. Compression of visual space before saccades. *Nature*. 1997; 386(6625):598–601.
- 960 **Rumelhart DE**, Hinton GE, Williams RJ. Learning representations by back-propagating errors. *nature*. 1986;
- 961 323(6088):533–536.
- 962 **Sagi D**, Julesz B. Enhanced detection in the aperture of focal attention during simple discrimination tasks.
- 963 *Nature*. 1986; 321(6071):693–695.
- 964 **Schrimpf M**, Kubilius J, Hong H, Majaj NJ, Rajalingham R, Issa EB, Kar K, Bashivan P, Prescott-Roy J, Schmidt K, et al. Brain-score: Which artificial neural network for object recognition is most brain-like? *BioRxiv*. 2018; p. 407007.
- 967 **Schwarzlose RF**, Swisher JD, Dang S, Kanwisher N. The distribution of category and location information across
- 968 object-selective regions in human visual cortex. *Proceedings of the National Academy of Sciences*. 2008;
- 969 105(11):4447–4452.
- 970 **Sclar G**, Maunsell JH, Lennie P. Coding of image contrast in central visual pathways of the macaque monkey.
- 971 *Vision research*. 1990; 30(1):1–10.
- 972 **Sprague TC**, Serences JT. Attention modulates spatial priority maps in the human occipital, parietal and frontal
- 973 cortices. *Nature neuroscience*. 2013; 16(12):1879–1887.
- 974 **Storrs KR**, Kietzmann TC, Walther A, Mehrer J, Kriegeskorte N. Diverse deep neural networks all predict human
- 975 IT well, after training and fitting. *bioRxiv*. 2020; .
- 976 **Theiss JD**, Bowen JD, Silver MA. Spatial Attention Enhances Crowded Stimulus Encoding Across Modeled Recep-
- 977 tive Fields by Increasing Redundancy of Feature Representations. *Neural computation*. 2022; 34(1):190–218.
- 978 **Tolias AS**, Moore T, Smirnakis SM, Tehovnik EJ, Siapas AG, Schiller PH. Eye movements modulate visual receptive
- 979 fields of V4 neurons. *Neuron*. 2001; 29(3):757–767.
- 980 **Treue S**, Trujillo JCM. Feature-based attention influences motion processing gain in macaque visual cortex.
- 981 *Nature*. 1999; 399(6736):575–579.
- 982 **Vo VA**, Sprague TC, Serences JT. Spatial tuning shifts increase the discriminability and fidelity of population
- 983 codes in visual cortex. *Journal of Neuroscience*. 2017; 37(12):3386–3401.
- 984 **Wagenmakers EJ**, Van Der Maas HL, Grasman RP. An EZ-diffusion model for response time and accuracy.
- 985 *Psychonomic bulletin & review*. 2007; 14(1):3–22.
- 986 **Wandell BA**, Winawer J. Imaging retinotopic maps in the human brain. *Vision research*. 2011; 51(7):718–737.
- 987 **Womelsdorf T**, Anton-Erxleben K, Pieper F, Treue S. Dynamic shifts of visual receptive fields in cortical area MT
- 988 by spatial attention. *Nature neuroscience*. 2006; 9(9):1156–1160.
- 989 **Yamins DL**, Hong H, Cadieu CF, Solomon EA, Seibert D, DiCarlo JJ. Performance-optimized hierarchical models
- 990 predict neural responses in higher visual cortex. *Proceedings of the national academy of sciences*. 2014;
- 991 111(23):8619–8624.
- 992 **Zirnsak M**, Steinmetz NA, Noudoost B, Xu KZ, Moore T. Visual space is compressed in prefrontal cortex before
- 993 eye movements. *Nature*. 2014 03; 507(7493):504 507. doi: 10.1038/nature13149.

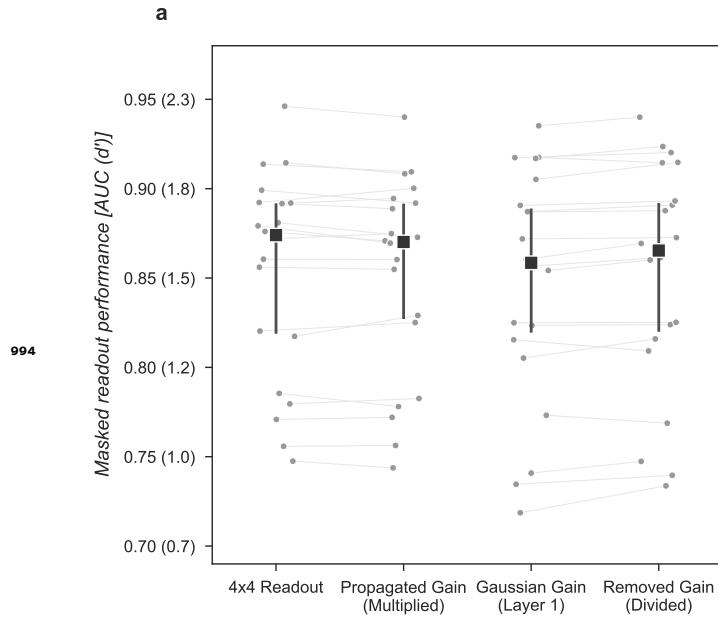


Figure 8—figure supplement 1. Direct readout from the cued quadrant improves performance alone, with no additional improvement from gain.

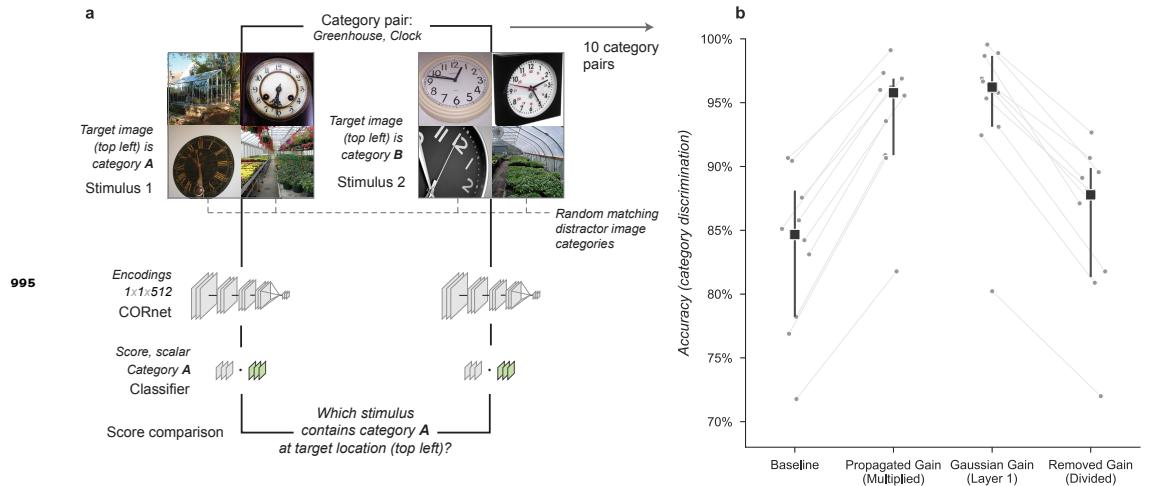


Figure 8—figure supplement 2. Gain propagation can account for changes in discrimination task performance due to Gaussian gain.